

# Hortonworks HDP on Amazon Web Services

## 1 INTRODUCTION

---

### 1.1 HORTONWORKS HDP

Hortonworks Data Platform (HDP) provides an enterprise ready data Hadoop platform that enables organizations to deploy the next generation Hadoop enterprise data platform.

Hortonworks Data Platform integrates with, and augments, your existing applications and systems so that you can take advantage of Hadoop with only minimal change to existing data architectures and skillsets. Deploy HDP in-cloud, on-premise or from an appliance across both Linux and Windows.

For more details, visit: <http://hortonworks.com/hdp/>

#### 1.1.1 Ambari

Ambari makes Hadoop management simpler by providing a consistent, secure platform for operational control. Ambari provides an intuitive Web UI as well as a robust REST API, which is particularly useful for automating cluster operations.

For more details, visit: <http://hortonworks.com/hadoop/ambari/>

### 1.2 AMAZON AWS

Amazon Web Services (AWS) provides on-demand computing resources and services in the cloud, with pay-as-you-go pricing.

For more details, visit: <http://docs.aws.amazon.com/gettingstarted/latest/awsgsg-intro/gsg-aws-intro.html>

#### 1.2.1 AWS EC2

Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides resizable compute capacity in the cloud. It is designed to make web-scale cloud computing easier for developers.

Benefits of EC2 services include:

- Elastic Web-Scale Computing
- Completely Controlled
- Flexible Cloud Hosting Services
- Designed for use with other Amazon Web Services
- Reliable, Secure and Inexpensive
- Easy to Start

For more details, visit: <https://aws.amazon.com/ec2/>

## 2 PREPARING YOUR INSTANCE

### 2.1 PREPARING VPC FOR INSTANCE AND CLUSTER

Amazon Virtual Private Cloud (Amazon VPC) lets you provision a logically isolated section of the Amazon Web Services (AWS) cloud where you can launch AWS resources in a virtual network that you define. You have complete control over your virtual networking environment, including selection of your own IP address range, creation of subnets, and configuration of route tables and network gateways.

For more details, visit: <https://aws.amazon.com/vpc/>

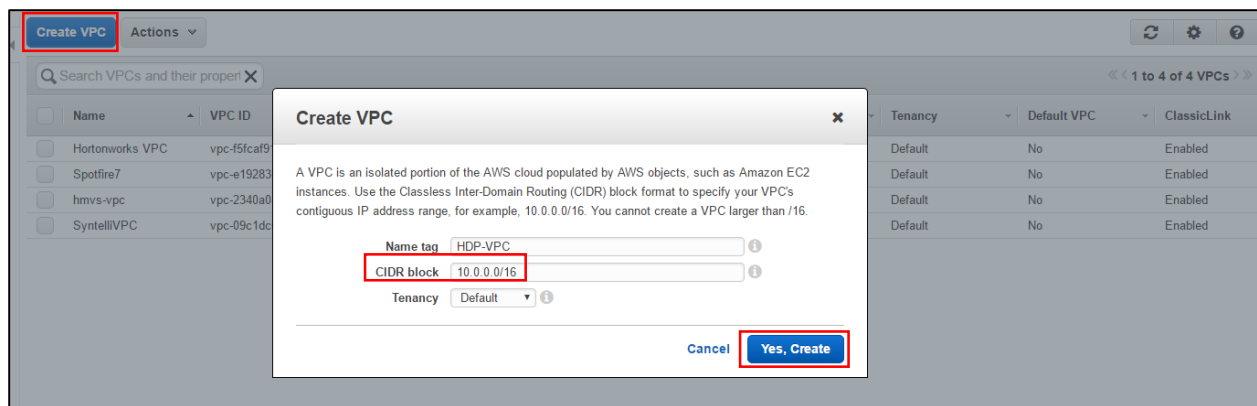


Figure 2.1.1

When I was creating a VPC, I had already created an AWS account. Once you have access to AWS Management Console, you can navigate to Networking. Under Networking, you will find VPC.

You can create a VPC using VPC Wizard or create one manually.

To create a VPC manually, follow the steps mentioned below:

1. Navigate to “Your VPCs” on the VPC dashboard.
2. Click “Create VPC” as highlighted, in Figure 2.1.1.
3. Provide a name for your VPC.
4. Next, provide a CIDR block. In my case, I entered (10.0.0.0/16)  
To learn more about CIDR, visit: [https://en.wikipedia.org/wiki/Classless\\_Inter-Domain\\_Routing](https://en.wikipedia.org/wiki/Classless_Inter-Domain_Routing)
5. Leave the Tenancy as “Default”. If choose “Dedicated”, then charges apply.  
To learn more about Dedicated, visit:  
<http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/dedicated-instance.html>
6. Finally click, “Yes, Create”.

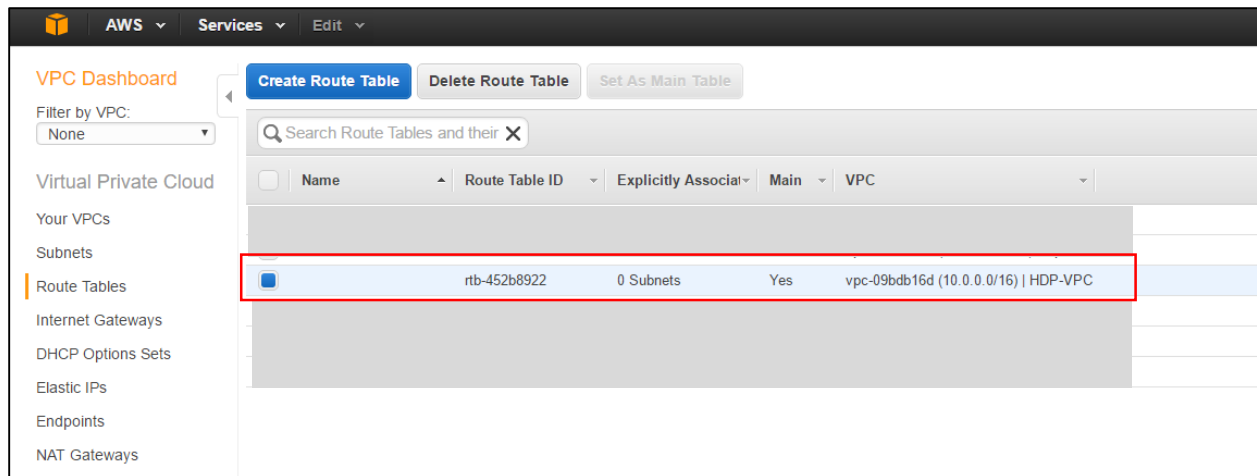


Figure 2.1.2

In Figure 2.1.2, if you navigate to Route Tables under VPC Dashboard, you will notice that a new Route Table is created along the VPC.




Figure 2.1.3

After creating a VPC, you can add one or more subnets in each Availability Zone. Each subnet must reside entirely within one Availability Zone and cannot span zones. Availability Zones are distinct locations that are engineered to be isolated from failures in other Availability Zones.

Navigate to Subnets under VPC Dashboard, and click “Create Subnet” as highlighted in Figure 2.1.3.

- Add a name-tag.
- Select the VPC to which the subnet will be linked to.
- Choose an availability zone, where the EC2 instances will be launched.

To read more about availability zone in AWS, visit:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

- Choose a CIDR block. In my case, I entered (10.0.1.0/24)
- Click “Yes, Create” as highlighted in Figure 2.1.4

After a successful subnet creation, you should be able to verify it, as shown in Figure 2.1.5

**Create Subnet**

Use the CIDR format to specify your subnet's IP address block (e.g., 10.0.0.0/24). Note that block sizes must be between a /16 netmask and /28 netmask. Also, note that a subnet can be the same size as your VPC.

Name tag: 10.0.1.0-us-east-1c

VPC: vpc-09bdbb16d (10.0.0.0/16) | HDP-VPC

Availability Zone: us-east-1c

CIDR block: 10.0.1.0/24

Value (us-east-1a) for parameter availabilityZone is invalid. Subnets can currently only be created in the following availability zones: us-east-1c, us-east-1b, us-east-1d, us-east-1e. (Service: AmazonEC2; Status Code: 400; Error Code: InvalidParameterValue; Request ID: 73ee663d-2401-4f2b-8cc8-a4b52a3b6f88)

Cancel Yes, Create

Figure 2.1.4

**VPC Dashboard**

Filter by VPC: None

Virtual Private Cloud

Your VPCs

**Subnets**

Route Tables

Internet Gateways

DHCP Options Sets

Elastic IPs

Endpoints

Create Subnet Subnet Actions

Q hdp

Name	Subnet ID	State	VPC	CIDR	Available IPs	Availability Zone
10.0.1.0-us-east-1c	subnet-99e68def	available	vpc-09bdbb16d (10.0.0.0/16)   HDP...	10.0.1.0/24	247	us-east-1c

Select a subnet above

Figure 2.1.5

An Internet gateway is a horizontally scaled, redundant, and highly available VPC component that allows communication between instances in your VPC and the Internet. It therefore imposes no availability risks or bandwidth constraints on your network traffic.

An Internet gateway serves two purposes:

- Provide a target in your VPC route tables for Internet-routable traffic
- Perform network address translation (NAT) for instances that have been assigned public IP addresses.

Navigate to Internet Gateways under VPC Dashboard and click “Create Internet Gateway” as highlighted in Figure 2.1.6.

Provide a name for the Internet Gateway and click “Yes, Create” as highlighted in Figure 2.1.7.

Choose VPC to which, the newly created Internet Gateway will be attached to. Next, click “Yes, Attach” as highlighted in Figure 2.1.6.

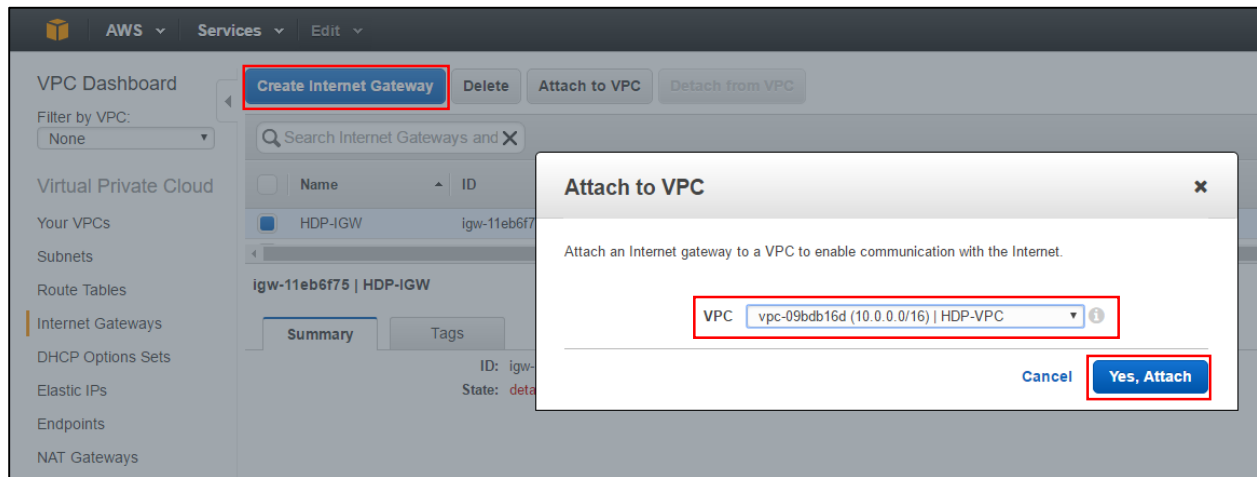


Figure 2.1.6

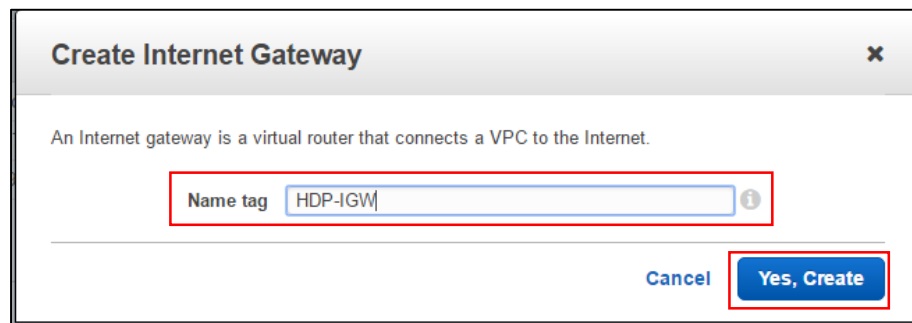


Figure 2.1.7

A route table contains a set of rules, called routes, that are used to determine where network traffic is directed.

Each subnet in your VPC must be associated with a route table; the table controls the routing for the subnet. A subnet can only be associated with one route table at a time, but you can associate multiple subnets with the same route table.

Navigate to Route Tables under VPC Dashboard and click “Create Route Table” as highlighted in Figure 2.1.8.

Provide a name for the Route Table and attach it to the VPC, that was created earlier. Click “Yes, Create”.

After the Route Table is created, you need to add an Internet route to it.

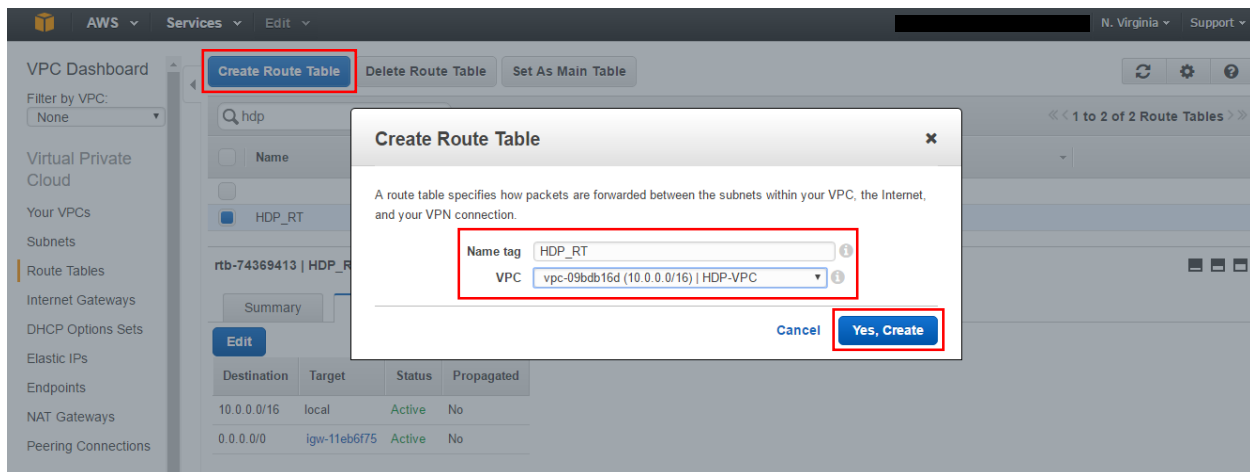


Figure 2.1.8

You can add an Internet Gateway route to your Route Table as highlighted in Figure 2.1.9. Select the route table and then select “Routes” tab as shown in Figure 2.1.9.

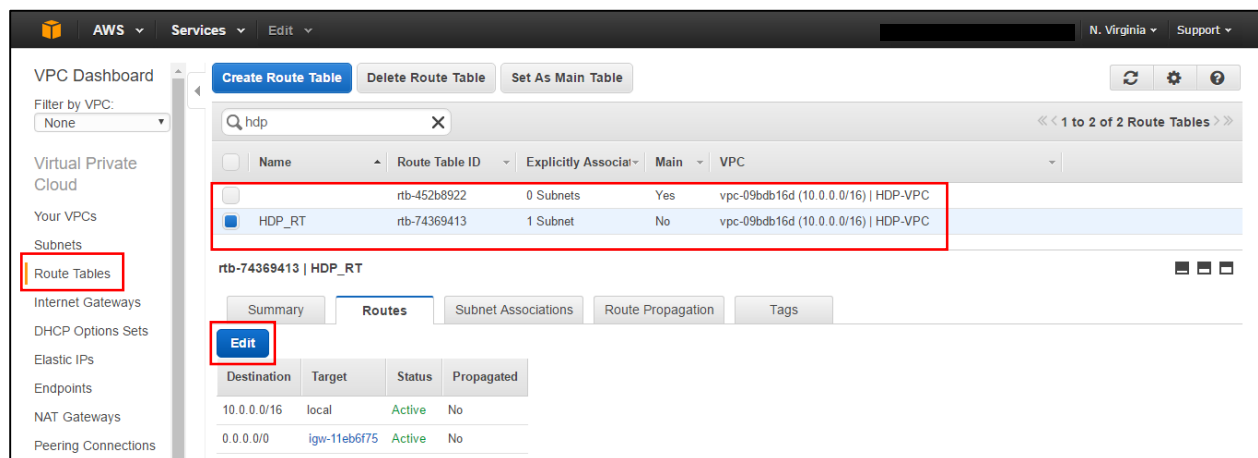


Figure 2.1.9

Click “Edit” and CIDR information. I entered (0.0.0.0/0) in the destination field, because I wanted the all the instances have Internet access as highlighted in Figure 2.1.10.

The Internet Gateway which was previous configured, was attached as a target. Click “Save”.

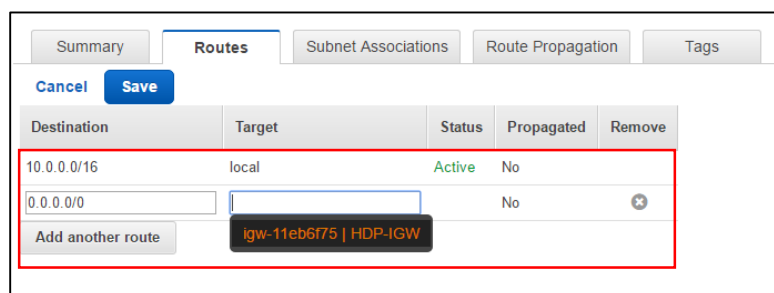


Figure 2.1.10




Figure 2.1.11

Next, Click “Subnet Association” tab and assign the subnet you created earlier to your Route Table. “Check” the subnet and click “Save” as highlighted in Figure 2.1.11

If you have associated the subnet to your Route Table, you will able to see it as shown in Figure 2.1.9, which states “1 Subnet” explicitly associated.

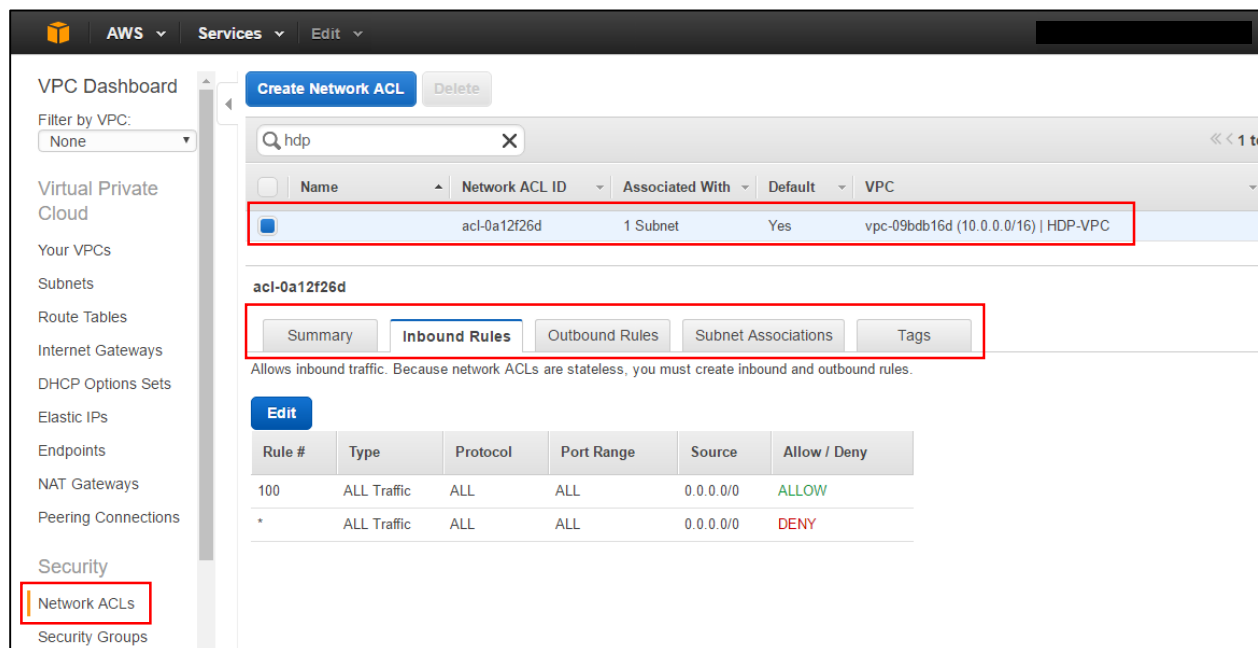


Figure 2.1.12

When a VPC is created, a Network ACL is also created for the VPC by default as highlighted in Figure 2.1.12. The Inbound and Outbound Rules are pre-defined. We can add rules if needed.

If you select the “Subnet Association” tab, you will notice that the subnet you created earlier is attached to the Network ACL.

To learn more about Network ACL, visit:

[http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC\\_ACLs.html](http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_ACLs.html)

Once you have completed all the steps to creating a VPC, you can proceed to launch an instance on AWS EC2.

## 2.2 PREPARING YOUR INSTANCE

Amazon EC2 provides a wide selection of instance types optimized to fit different use cases.

Instance types comprise varying combinations of CPU, memory, storage, and networking capacity and give you the flexibility to choose the appropriate mix of resources for your applications. Each instance type includes one or more instance sizes, allowing you to scale your resources to the requirements of your target workload.

I chose M3 instance type. M3 instance types and provides a balance of compute, memory, and network resources, and it is a good choice for many applications.

For more details about EC2 instance types, visit: <https://aws.amazon.com/ec2/instance-types/>

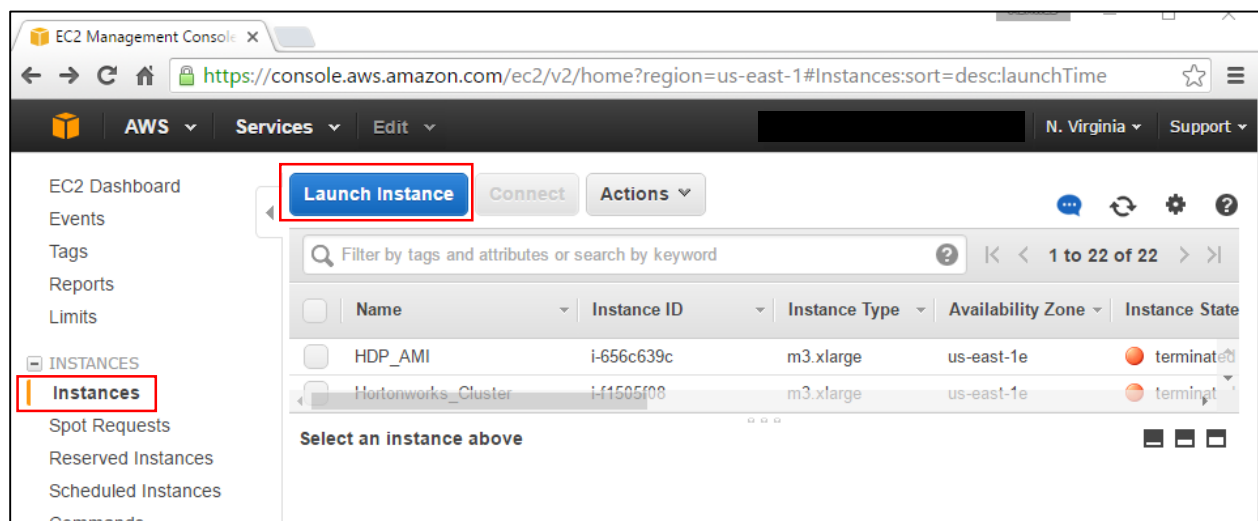


Figure 2.2.1

Navigate to EC2 Management Console from AWS Home screen.

Under EC2 Dashboard, click “Instances” and click “Launch Instance” as highlighted in Figure 2.2.1.

The next screen is shown Figure 2.2.2. We need to choose an Amazon Machine Image.

An Amazon Machine Image (AMI) provides the information required to launch an instance, which is a virtual server in the cloud. You specify an AMI when you launch an instance, and you can launch as many instances from the AMI as you need.

An AMI includes the following:

- A template for the root volume for the instance (for example, an operating system, an application server, and applications)
- Launch permissions that control which AWS accounts can use the AMI to launch instances.
- A block device mapping that specifies the volumes to attach to the instance when it's launched



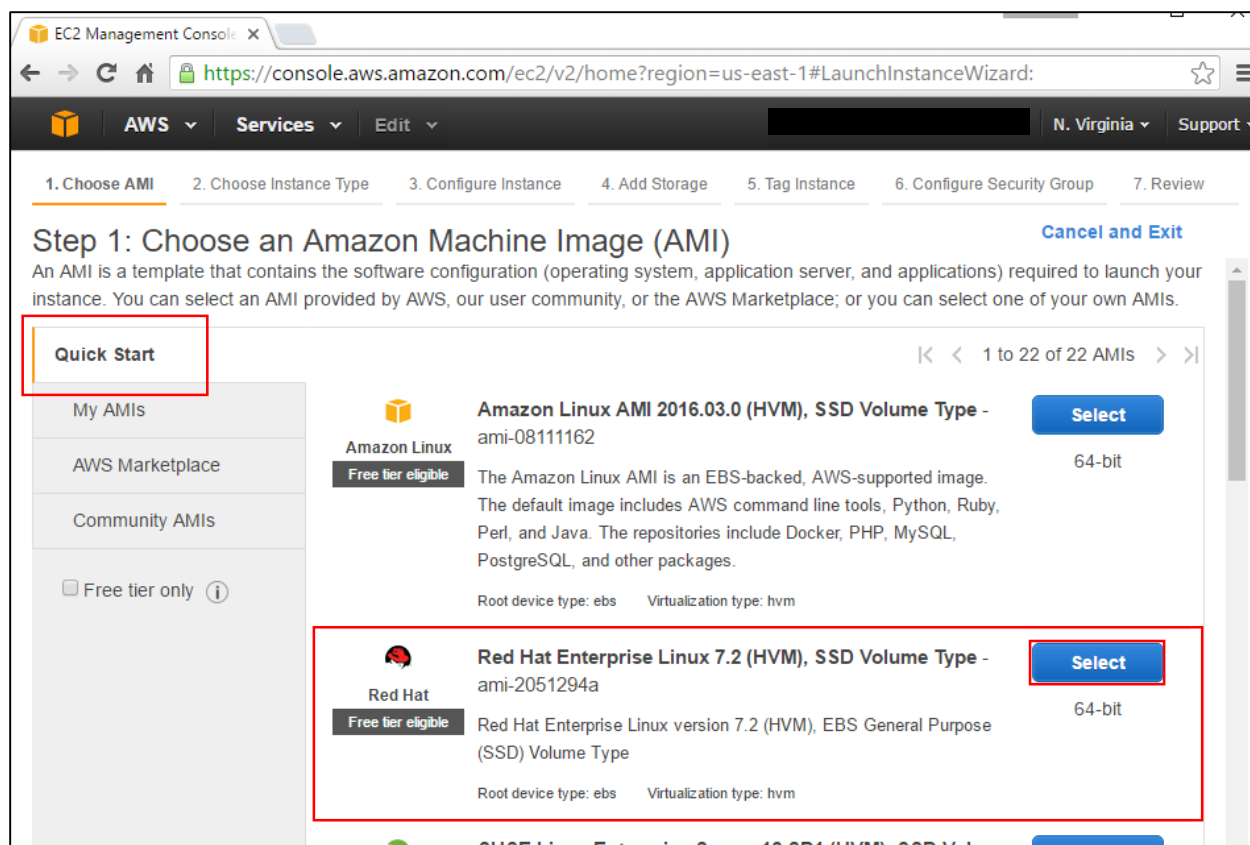


Figure 2.2.2

I chose Red Hat Enterprise Linux (RHEL) 7.2 as highlighted in Figure 2.2.2. Most of the generic AMIs are present in “Quick Start”. Click “Select”

The next screen will ask you to choose an instance type as shown in Figure 2.2.3.

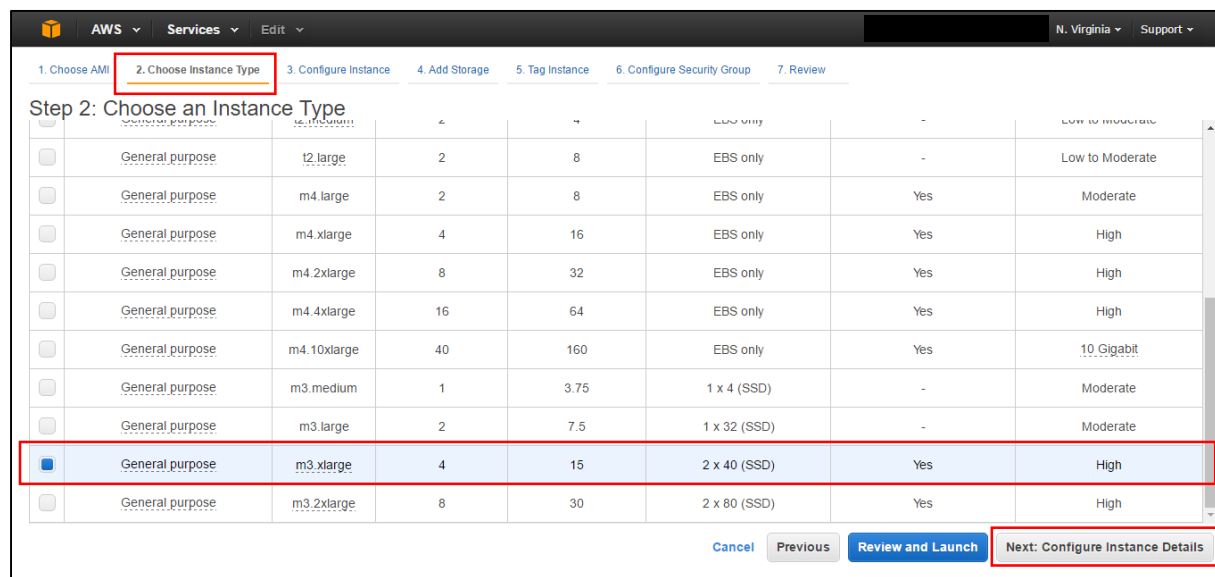


Figure 2.2.3

I chose “m3.xlarge” with 4 virtual CPUs, 15 GB of RAM and an instance storage of 2 X 40 GB SSD (EBS Optimized). The instance also provides high network performance.

Click “Next: Configure Instance Details” to proceed.

The screenshot shows the AWS Management Console interface for configuring an instance. The top navigation bar includes the AWS logo, 'Services', and 'Edit'. Below the navigation bar is a progress bar with steps: 1. Choose AMI, 2. Choose Instance Type, 3. Configure Instance Details (active), 4. Add Storage, 5. Tag Instance, 6. Configure Security Group, and 7. Review. The main heading is 'Step 3: Configure Instance Details'. Below the heading is a subheading: 'Configure the instance to suit your requirements. You can launch multiple instances from the same AMI, request Spot instances to take advantage of the lower pricing, assign an access key, and more.' The configuration options are as follows:

- Number of instances:** 1 (highlighted with a red box). A link 'Launch into Auto Scaling Group' is next to it.
- Purchasing option:** ☐ Request Spot instances.
- Network:** vpc-09bdbb16d (10.0.0.0/16) | HDP-VPC (highlighted with a red box). A link 'Create new VPC' is next to it.
- Subnet:** subnet-99e68def (10.0.1.0/24) | 10.0.1.0-us-east-1c (highlighted with a red box). A link 'Create new subnet' is next to it.
- Auto-assign Public IP:** Enable (highlighted with a red box).
- IAM role:** None (highlighted with a red box). A link 'Create new IAM role' is next to it.
- Shutdown behavior:** Stop.
- Enable termination protection:** ☐ Protect against accidental termination.
- Monitoring:** ☐ Enable CloudWatch detailed monitoring. A link 'Additional charges apply.' is next to it.
- EBS-optimized instance:** ☐ Launch as EBS-optimized instance. A link 'Additional charges apply.' is next to it.
- Tenancy:** Shared - Run a shared hardware instance. A link 'Additional charges will apply for dedicated tenancy.' is next to it.

Figure 2.2.4

In Configure Instance Details, we specify the number of instances, type of network and subnet we want to launch. Shown in Figure 2.2.4

I chose one instance, since I need to configure the instance for launching HDP on AWS.

For network, I selected the VPC that I had created in the earlier section.

I attached the subnet, which I created in the previous section. Also, choose “Enable” for “Auto-assign Public IP”. This will assign a Public IP, which will be used to SSH into the instance.

Next step in launching an instance is adding storage. The instance selected has a default volume 10 GB, which is set for “Linux Root”, highlighted in yellow.

Click “Add New Volume”.

A new list populates. Select “EBS” as the volume type. The device, is the path to the disk space as highlighted in Figure 2.2.5. I selected my path to “/dev/sdb”. I entered my volume size as 50 GB.

To learn more about AWS EBS, visit: <https://aws.amazon.com/ebs/>

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Tag Instance 6. Configure Security Group 7. Review

### Step 4: Add Storage

Your instance will be launched with the following storage device settings. You can attach additional EBS volumes and instance store volumes to your instance, or edit the settings of the root volume. You can also attach additional EBS volumes after launching an instance, but not instance store volumes. [Learn more](#) about storage options in Amazon EC2.

Volume Type ⓘ	Device ⓘ	Snapshot ⓘ	Size (GiB) ⓘ	Volume Type ⓘ	IOPS ⓘ	Delete on Termination ⓘ	Encrypted ⓘ
Root	/dev/sda1	snap-ba40cac8	10	General Purpose ⚙	30 / 3000	<input checked="" type="checkbox"/>	Not Encrypted
EBS	/dev/sdb	Search (case-insensit	50	General Purpose ⚙	150 / 3000	<input type="checkbox"/>	<input type="checkbox"/>

[Add New Volume](#)

Free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage. [Learn more](#) about free usage tier eligibility and usage restrictions.

[Cancel](#) [Previous](#) [Review and Launch](#) [Next: Tag Instance](#)

Figure 2.2.5

Next step is to tag your instance. Click “Next: Configure Security Group”

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Tag Instance 6. Configure Security Group 7. Review

### Step 5: Tag Instance

A tag consists of a case-sensitive key-value pair. For example, you could define a tag with key = Name and value = Webserver. [Learn more](#) about tagging your Amazon EC2 resources.

Key (127 characters maximum)	Value (255 characters maximum)
Name	HDP_AMI
Owner	Vishwas

[Create Tag](#) (Up to 10 tags maximum)

[Cancel](#) [Previous](#) [Review and Launch](#) [Next: Configure Security Group](#)

Figure 2.2.6

**Step 6: Configure Security Group**

A security group is a set of firewall rules that control the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web server and allow Internet traffic to reach your instance, add rules that allow unrestricted access to the HTTP and HTTPS ports. You can create a new security group or select from an existing one below. [Learn more](#) about Amazon EC2 security groups.

**Assign a security group:** ☒ Create a new security group ☐ Select an existing security group

**Security group name:**

**Description:**

Type	Protocol	Port Range	Source
SSH	TCP	22	Anywhere 0.0.0.0/0
All TCP	TCP	0 - 65535	Anywhere 0.0.0.0/0
All UDP	UDP	0 - 65535	Anywhere 0.0.0.0/0
All ICMP	ICMP	0 - 65535	Anywhere 0.0.0.0/0
All traffic	All	0 - 65535	Anywhere 0.0.0.0/0

**Add Rule**

**Warning**  
Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

[Cancel](#) [Previous](#) [Review and Launch](#)

Figure 2.2.7

In this step, you will need to create a new security group or choose an existing security group.

A security group acts as a virtual firewall that controls the traffic for one or more instances. When you launch an instance, you associate one or more security groups with the instance.

To learn more about Security Groups in AWS, visit:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-network-security.html>

Provide a name for your security group. You need to add rules, which will allow you to access your instance using several protocols.

Source defines the user access to the instance. Setting the source (0.0.0.0/0) makes your instance accessible from any device, when the obtain the Public IP address of the instance. The above link about security group explains, how your instance access can be made more secure.

For Proof of Concept (POC), the above configurations shown in Figure 2.2.7 will be helpful.

Next, click “Review and Launch”.

1. Choose AMI
2. Choose Instance Type
3. Configure Instance
4. Add Storage
5. Tag Instance
6. Configure Security Group
7. Review

## Step 7: Review Instance Launch

Please review your instance launch details. You can go back to edit changes for each section. Click **Launch** to assign a key pair to your instance and complete the launch process.

**Improve your instances' security. Your security group, HDP\_SG, is open to the world.**  
Your instances may be accessible from any IP address. We recommend that you update your security group rules to allow access from known IP addresses only.  
You can also open additional ports in your security group to facilitate access to the application or service you're running, e.g., HTTP (80) for web servers. [Edit security groups](#)

**Your instance configuration is not eligible for the free usage tier**

To launch an instance that's eligible for the free usage tier, check your AMI selection, instance type, configuration options, or storage devices. Learn more about [free usage tier](#) eligibility and usage restrictions.

[Don't show me this again](#)

▼
AMI Details
Edit AMI

**Red Hat Enterprise Linux 7.2 (HVM), SSD Volume Type - ami-2051294a**  

Free tier eligible
Red Hat Enterprise Linux version 7.2 (HVM), EBS General Purpose (SSD) Volume Type  
Root Device Type: ebs    Virtualization type: hvm

▼
Instance Type
Edit instance type

Instance Type	ECUs	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance
m3.medium	3	1	3.75	1 x 4	-	Moderate

▼
Security Groups
Edit security groups

Security group name: HDP\_SG  
Description: HDP\_SG created 2016-03-27T18:03:25.483-04:00

Cancel
Previous
Launch

Figure 2.2.8

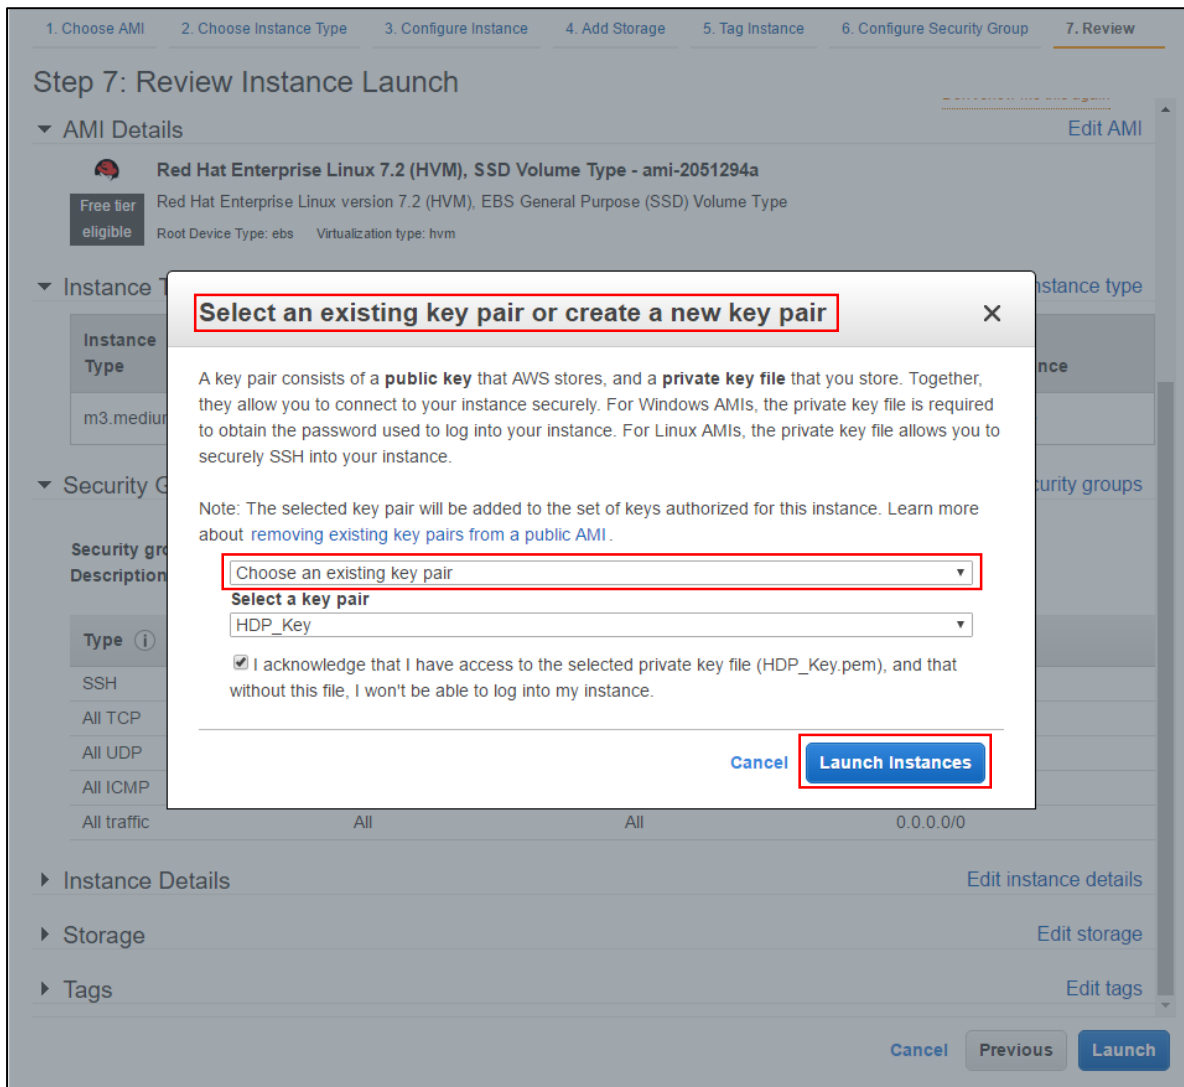


Figure 2.2.9

As you approach, the last step to launch your instance, you will be asked to select a key-pair. You can either create a new one, or select an existing key-pair as highlighted in Figure 2.2.9.

Launch Instance Connect Actions

Filter by tags and attributes or search by keyword

1 to 23 of 23

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks
<input checked="" type="checkbox"/> HDP_AMI	i-e66a651f	m3.medium	us-east-1e	running	2/2 checks ...
<input type="checkbox"/> HDP_AMI	i-656c639c	m3.xlarge	us-east-1e	terminated	

Instance: **i-e66a651f (HDP\_AMI)** Public DNS: **ec2-54-165-143-171.compute-1.amazonaws.com**

Description Status Checks Monitoring Tags

Instance ID	i-e66a651f	Public DNS	ec2-54-165-143-171.compute-1.amazonaws.com
Instance state	running	Public IP	<b>54.165.143.171</b>
Instance type	m3.medium	Elastic IP	-
Private DNS	<b>ip-10-0-0-40.ec2.internal</b>	Availability zone	us-east-1e
Private IPs	<b>10.0.0.40</b>	Security groups	HDP_SG. <a href="#">view rules</a>
Secondary private IPs		Scheduled events	No scheduled events
VPC ID	vpc-2340a047	AMI ID	RHEL-7.2_HVM_GA-20151112-x86_64-1-Hourly2-GP2 (ami-2051294a)

Figure 2.2.10

Figure 2.2.10 highlights the instance has launched successfully and is running.

## 3 CONFIGURING YOUR INSTANCE FOR HDP INSTALLATION

---

### 3.1 INSTALL APACHE




Figure 3.1.1

Once your instance is ready and running as shown in Figure 2.2.10 (previous section), then you should be able to log into the machine by SSH. Login as “ec2-user”.

To learn more about connecting to Linux instances on AWS, visit:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AccessingInstances.html>

Just to make sure, the instance public IP address is visible, install Apache.

To install Apache, follow instructions in the link below:

<http://www.cyberciti.biz/faq/howto-install-linux-apache-mariadb-php-lamp-stack-on-centos7-rhel7/>

If Apache is installed successfully, you should see the following webpage online accessible on your instance. Shown in Figure 3.1.2.



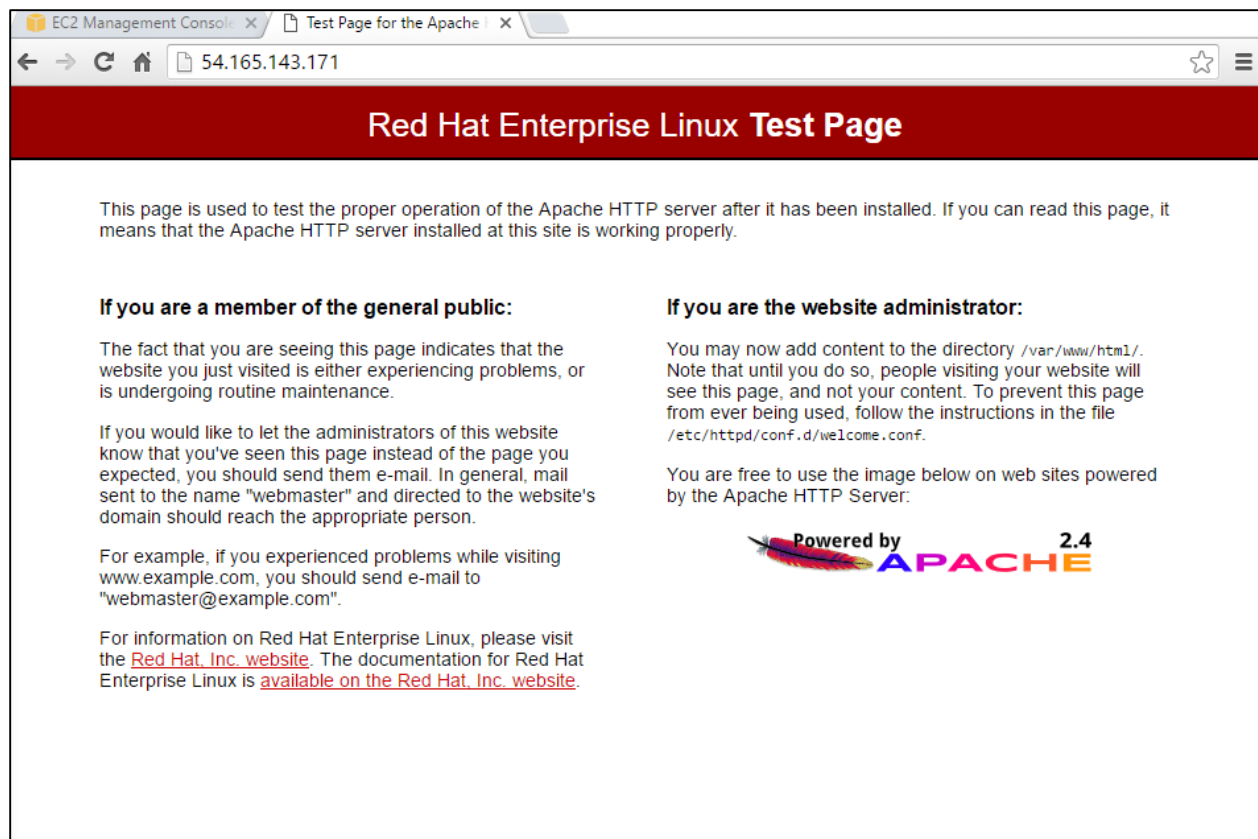


Figure 3.1.2

## 3.2 ATTACH ADDITIONAL VOLUME (EBS) TO YOUR INSTANCE

Remember, we attached an EBS of 50 GB for our instance. Now, we need to verify, if this storage is available.

```
$ df -h
```

The **df** command provides an option to display sizes in Human Readable formats by using **'-h'** (prints the results in human readable format), highlighted in Figure 3.2.1.

To learn more about disk file system, visit: <http://www.tecmint.com/how-to-check-disk-space-in-linux/>

```
$ lsblk
```

The **lsblk** command allows you to display a list of available block devices. Highlighted in Figure 3.2.1.

For each listed block device, the **lsblk** command displays the device name (NAME), major and minor device number (MAJ:MIN), if the device is removable (RM), what is its size (SIZE), if the device is read-only (RO), what type is it (TYPE), and where the device is mounted (MOUNTPOINT).

```
ec2-user@ip-10-0-0-40:~$ df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/xvda2       10G   1.1G   9.0G  11% /
devtmpfs         1.8G    0   1.8G    0% /dev
tmpfs            1.8G    0   1.8G    0% /dev/shm
tmpfs            1.8G   17M   1.8G    1% /run
tmpfs            1.8G    0   1.8G    0% /sys/fs/cgroup
tmpfs            361M    0   361M    0% /run/user/1000
[ec2-user@ip-10-0-0-40 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   10G  0 disk
├─xvda1     202:1    0    1M  0 part
└─xvda2     202:2    0   10G  0 part /
xvdb        202:16   0    50G  0 disk
[ec2-user@ip-10-0-0-40 ~]$ sudo mkfs -t ext4 /dev/xvdb
mke2fs 1.42.9 (28-Dec-2013)
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=0 blocks
3276800 inodes, 13107200 blocks
655360 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=2162163712
400 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624, 11239424

Allocating group tables: done
Writing inode tables: done
Creating journal (32768 blocks): done
Writing superblocks and filesystem accounting information: done

[ec2-user@ip-10-0-0-40 ~]$ sudo mkdir /grid
[ec2-user@ip-10-0-0-40 ~]$ sudo mount /dev/xvdb /grid
[ec2-user@ip-10-0-0-40 ~]$
```

Figure 3.2.1

To learn more about `lsblk`, visit:

[https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/6/html/Deployment\\_Guide/s1-sysinfo-filesystems.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Deployment_Guide/s1-sysinfo-filesystems.html)

The `lsblk` command, displays the attached 50 GB EBS as “xvdb”, highlighted in red, in Figure 3.2.1.

Once you add a new disk drive to the system, you may want to partition the drive and use the **ext4** file system.

```
$ sudo mkfs -t ext4 /dev/xvdb
```

Highlighted in yellow, in Figure 3.2.1.

To learn more about make file system, refer:

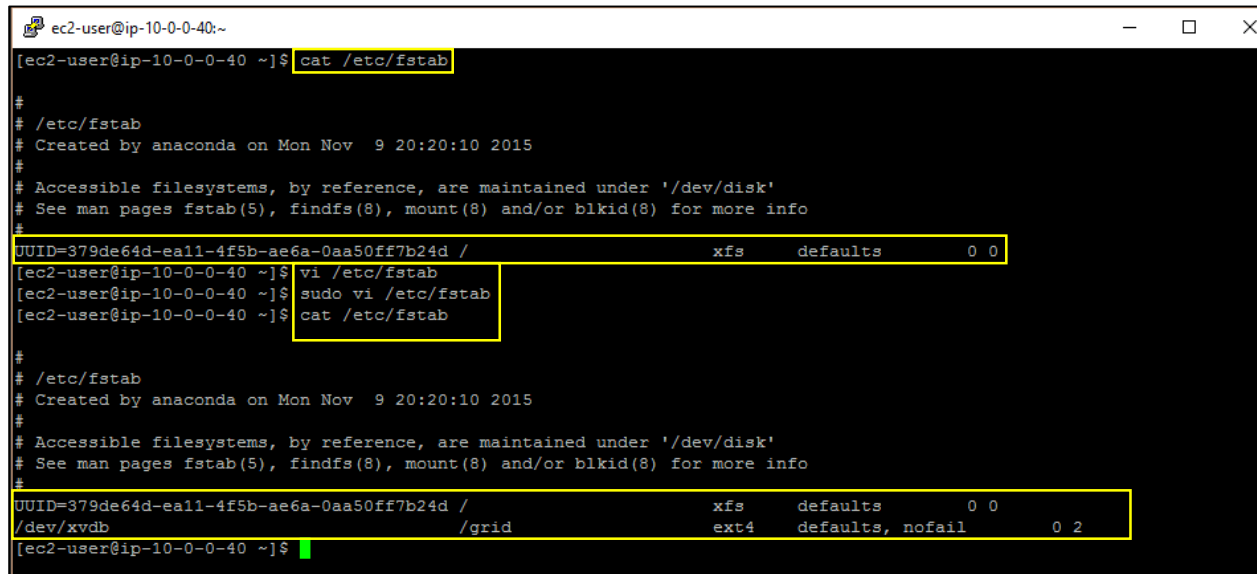
[https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/5/html/Deployment\\_Guide/s1-file-system-ext4-create.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/5/html/Deployment_Guide/s1-file-system-ext4-create.html)

Next, we need to create a directory “grid” in the **root directory**.

```
$ sudo mkdir /grid
```

```
$ sudo mount /dev/xvdb /grid
```

We can now mount the new storage to “grid” directory on the AWS EC2 instance.



```
ec2-user@ip-10-0-0-40:~$ cat /etc/fstab
#
# /etc/fstab
# Created by anaconda on Mon Nov  9 20:20:10 2015
#
# Accessible filesystems, by reference, are maintained under '/dev/disk'
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info
#
UUID=379de64d-ea11-4f5b-ae6a-0aa50ff7b24d / xfs defaults 0 0
[ec2-user@ip-10-0-0-40 ~]$ vi /etc/fstab
[ec2-user@ip-10-0-0-40 ~]$ sudo vi /etc/fstab
[ec2-user@ip-10-0-0-40 ~]$ cat /etc/fstab
#
# /etc/fstab
# Created by anaconda on Mon Nov  9 20:20:10 2015
#
# Accessible filesystems, by reference, are maintained under '/dev/disk'
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info
#
UUID=379de64d-ea11-4f5b-ae6a-0aa50ff7b24d / xfs defaults 0 0
/dev/xvdb /grid ext4 defaults, nofail 0 2
[ec2-user@ip-10-0-0-40 ~]$
```

Figure 3.2.2

The **/etc/fstab** file is used to control what file systems are mounted when the system boots, as well as to supply default values for other file systems that may be mounted manually from time to time.

As highlighted in Figure 3.2.2, I listed the contents present in “**fstab**” file.

Next, you need to edit the “**fstab**” file and add content as highlighted in Figure 3.4.

/dev/xvdb	/grid	ext4	defaults, nofail	0 2
-----------	-------	------	------------------	-----

To learn more about **fstab**, visit:

[https://access.redhat.com/documentation/en-US/Red Hat Enterprise Linux/4/html/Introduction\\_To\\_System\\_Administration/s2-storage-mount-fstab.html](https://access.redhat.com/documentation/en-US/Red Hat Enterprise Linux/4/html/Introduction_To_System_Administration/s2-storage-mount-fstab.html)

Finally **reboot** your instance.

### 3.3 DISABLE SELINUX AND ENABLE NTP

Security-Enhanced Linux (**SELinux**) is an implementation of a mandatory access control mechanism in the Linux kernel, checking for allowed operations after standard discretionary access controls are checked.

SELinux can enforce rules on files and processes in a Linux system, and on their actions, based on defined policies.

To learn more about SELinux, visit:

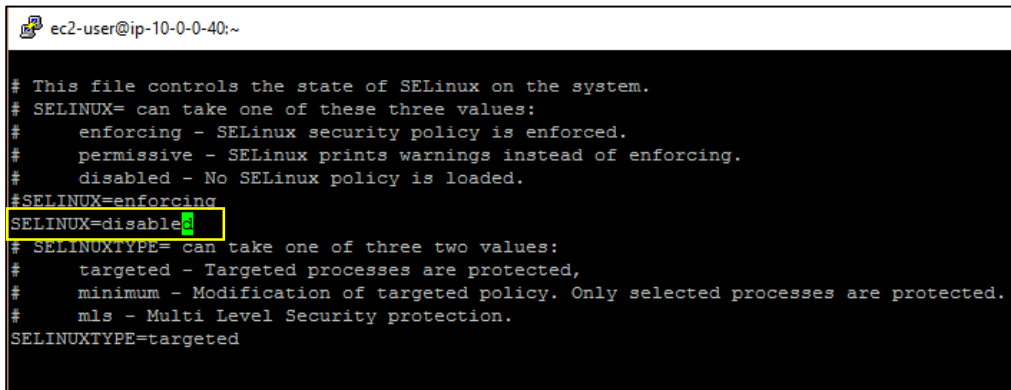
[https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/7/html/SELinux\\_Users\\_and\\_Administrators\\_Guide/chap-Security-Enhanced\\_Linux-Introduction.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/SELinux_Users_and_Administrators_Guide/chap-Security-Enhanced_Linux-Introduction.html)

When SELinux is disabled, SELinux policy is not loaded at all; it is not enforced.

To permanently disable SELinux, follow the procedure below:

```
$ sudo vi /etc/sysconfig/selinux
```

Configure **SELINUX=disabled** in the **/etc/selinux/config** file, highlighted in Figure 3.3.1.



```
ec2-user@ip-10-0-0-40:~  
# This file controls the state of SELinux on the system.  
# SELINUX= can take one of these three values:  
#   enforcing - SELinux security policy is enforced.  
#   permissive - SELinux prints warnings instead of enforcing.  
#   disabled - No SELinux policy is loaded.  
#SELINUX=enforcing  
SELINUX=disabled  
# SELINUXTYPE= can take one of three two values:  
#   targeted - Targeted processes are protected,  
#   minimum - Modification of targeted policy. Only selected processes are protected.  
#   mls - Multi Level Security protection.  
SELINUXTYPE=targeted
```

Figure 3.3.1

Next, you need to disable “**firewall daemon**” in your instance.

```
$ sudo yum install firewalld
```

If the instance, throws a message stating the service is unavailable, then you need to install firewall service.

```
$ sudo yum install firewalld
```

After installation completes, you can disable “**firewalld**” using “**sudo yum install firewalld**”.

Next, stop the “**firewalld**” service.

```
$ sudo service firewalld stop
```

NTP (**N**etwork **T**ime **P**rotocol) is a protocol to keep servers time synchronized: one or several master servers provide time to client servers that can themselves provide time to other client servers (notion of stratus).

You need NTP service on your instance. You need install and enable NTP on start-up.

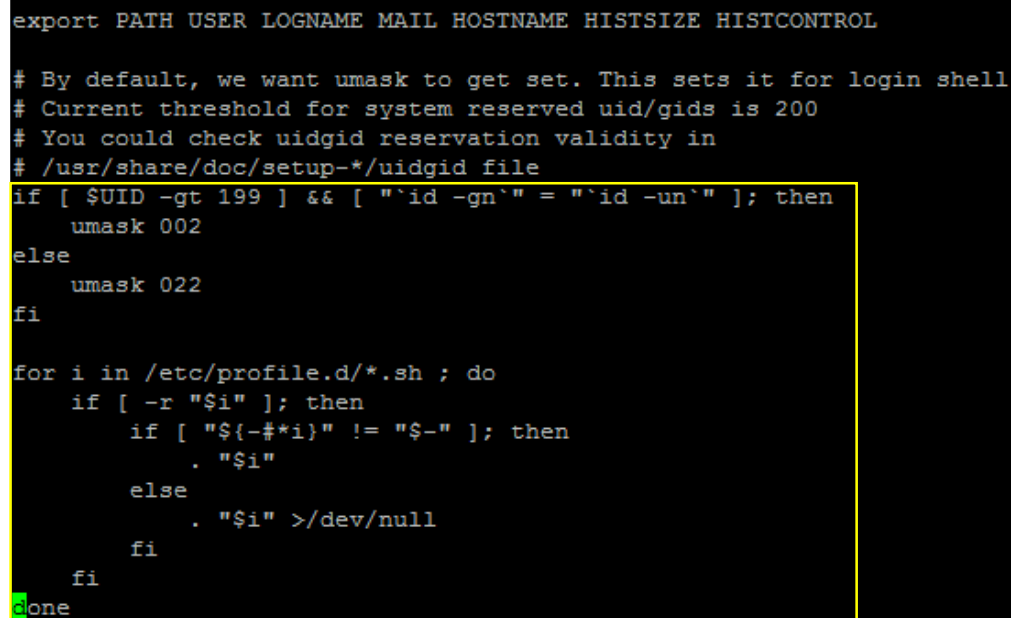
```
$ sudo yum install ntp
$ sudo systemctl enable ntpd
$ sudo systemctl start ntpd
$ sudo systemctl is-enabled ntpd
```

### 3.4 UMASK

UMASK (**U**ser **M**ask or User file creation MASK) is the default permission or base permissions given when a new file (even folder too, as Linux treats everything as files) is created on a Linux machine. Most of the Linux distros give 022 (**0022**) as default UMASK. In other words, it is a system default permission for newly created files/folders in the machine.

Modify the “/etc/profile” file. Remove the “**umask**” condition as highlighted in Figure 3.4.1 and set it to “**umask 022**” as highlighted in Figure 3.4.2.

```
$ umask
$ sudo vi /etc/profile
```



```
export PATH USER LOGNAME MAIL HOSTNAME HISTSIZE HISTCONTROL

# By default, we want umask to get set. This sets it for login shell
# Current threshold for system reserved uid/gids is 200
# You could check uidgid reservation validity in
# /usr/share/doc/setup-*/uidgid file
if [ $UID -gt 199 ] && [ "`id -gn`" = "`id -un`" ]; then
    umask 002
else
    umask 022
fi

for i in /etc/profile.d/*.sh ; do
    if [ -r "$i" ]; then
        if [ "${-#*i}" != "$-" ]; then
            . "$i"
        else
            . "$i" >/dev/null
        fi
    fi
done
```

Figure 3.4.1

```

export PATH USER LOGNAME MAIL HOSTNAME HISTSIZE HISTCONTROL

# By default, we want umask to get set. This sets it for login shell
# Current threshold for system reserved uid/gids is 200
# You could check uidgid reservation validity in
# /usr/share/doc/setup-*/uidgid file
#if [ $UID -gt 199 ] && [ "`id -gn`" = "`id -un`" ]; then
#    umask 002
#else
umask 022
#fi

```

Figure 3.4.2

As seen in Figure 3.4.2, the red-highlighted boxes are commented out. The yellow box, shows the “**umask**”, set to “**umask 022**”.

### 3.5 CREATE SSH KEYGEN

**SSH** (Secure Shell) is a protocol which facilitates secure communications between two systems using a client-server architecture and allows users to log in to server host systems remotely. Unlike other remote communication protocols, such as FTP or Telnet, SSH encrypts the login session, rendering the connection difficult for intruders to collect unencrypted passwords.

For SSH to be truly effective, using insecure connection protocols should be prohibited. Otherwise, a user's password may be protected using SSH for one session, only to be captured later while logging in using Telnet.

```
$ ssh-keygen -t rsa
```

```

ec2-user@ip-10-0-0-40:~$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/ec2-user/.ssh/id_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/ec2-user/.ssh/id_rsa.
Your public key has been saved in /home/ec2-user/.ssh/id_rsa.pub.
The key fingerprint is:
3b:fa:05:a7:d9:10:59:c4:84:a0:e8:ed:6c:92:c5:af ec2-user@ip-10-0-0-40.ec2.internal
The key's randomart image is:
+--[ RSA 2048 ]-----+
|      .. =+      |
|      . . .o.    |
|      . . o      |
|      . o .      |
|      . + S .    |
|      = . O      |
|      o + . = o   |
|      o . . o     |
|      E ...      |
+-----+
[ec2-user@ip-10-0-0-40 ~]$

```

Figure 3.5.1

Reboot the instance to apply the changes made to the instance.

### 3.6 VERIFY UMASK AND NTP

Run the following command mentioned below to verify, if UMASK and NTP are enabled.

```
$ umask  
$ sudo systemctl is-enabled ntpd
```

Next, move the public key to “authorized\_keys” file.

The public key was created during “Create SSH Keygen” step.

```
$ ls .ssh/  
$ cd .ssh  
$ $ cat id_rsa.pub >> authorized_keys
```

To verify, a password-less authentication is successful, SSH into the same instance using your instance’s

Private IP or Private DNS

For example:

My private IP is **10.0.0.40**

My private DNS is **ip-10-0-0-40.ec2.internal**

```
$ ssh ip-10-0-0-40.ec2.internal
```

When you enter the above command, you should login, without entering any passwords.

If your login was successful, that means, you have configured password-less SSH correctly.

To read more about password-less SSH, visit:

[https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/7/html/System\\_Administrators\\_Guide/s1-ssh-configuration.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/System_Administrators_Guide/s1-ssh-configuration.html)

## 4 CREATE A BASE IMAGE ON AWS MANAGEMENT CONSOLE

After your instance has been configured for installing HDP, we need to make a base image.

This image will be later utilized to create a HDP cluster.

The base image helps you to replicate same configuration on several host machines, avoiding all the steps performed in Section 2.

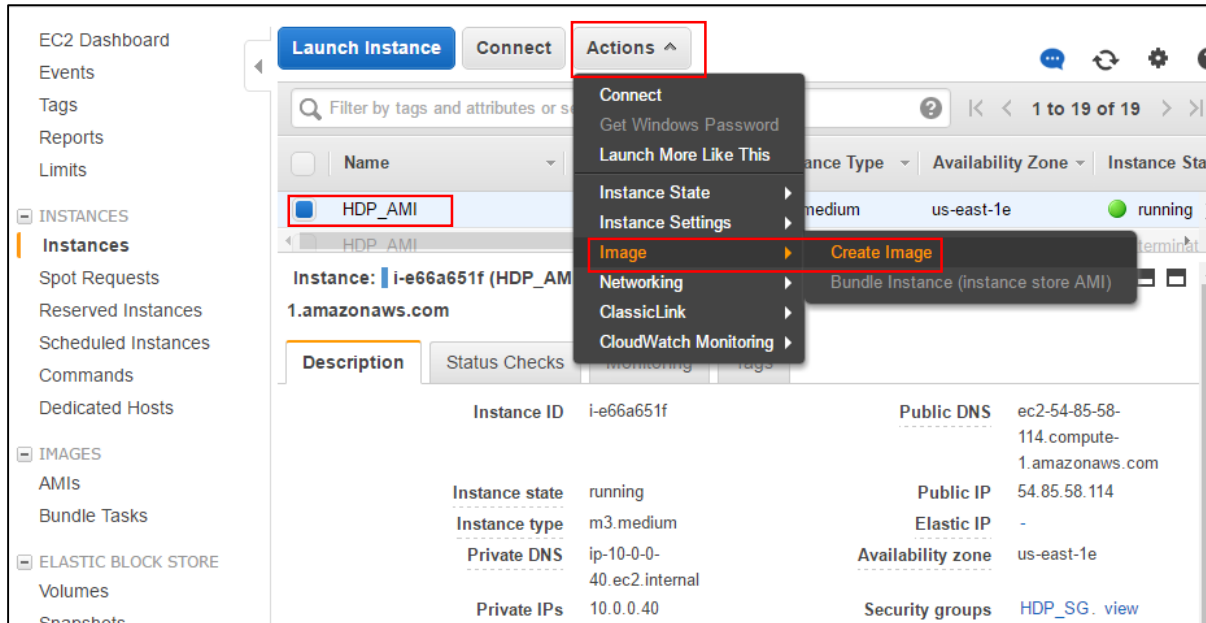


Figure 4.1

Navigate to Instances on EC2 Dashboard and select the configured instance as highlighted in Figure 4.1.

Click "Actions" and navigate to "Image". Then click on "Create Image".

You will get an alert dialog box, saying "Create image request received" as highlighted in Figure 4.2

Once the base image has been created, navigate to AMIs under EC2 Dashboard as highlighted in Figure 4.3. You will notice, that a new AMI is created and available for use.



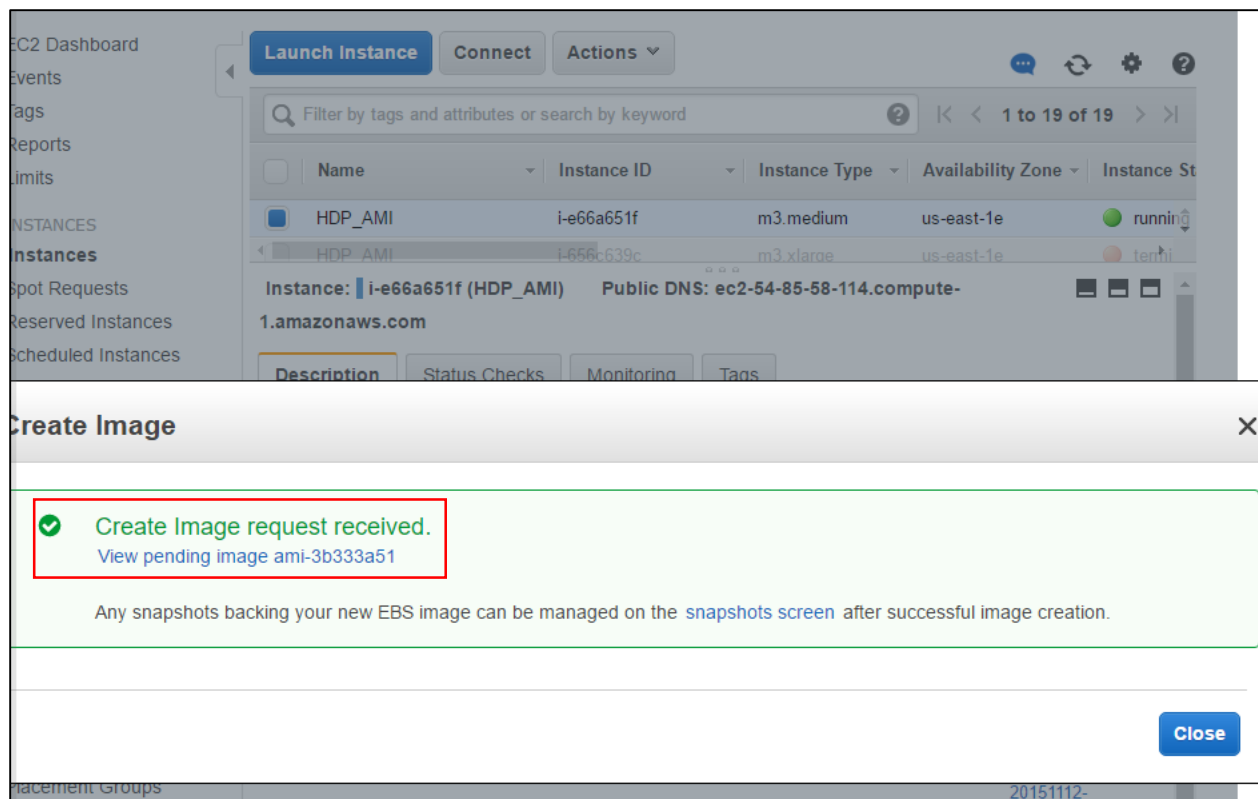


Figure 4.2

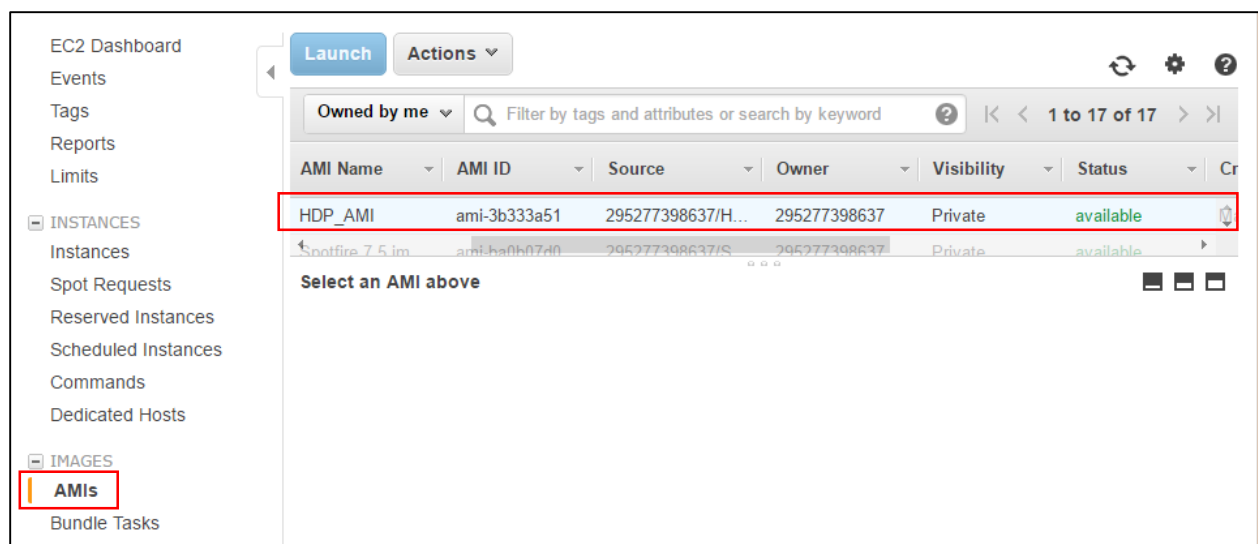


Figure 4.3

## 5 CONFIGURE CLUSTER USING BASE IMAGE

We will create a HDP cluster, with the help of Base Image created earlier in Section 4.

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Tag Instance 6. Configure Security Group 7. Review

### Step 3: Configure Instance Details

Configure the instance to suit your requirements. You can launch multiple instances from the same AMI, request Spot instances to take advantage of the low cost of Spot instances.

**Number of instances** ⓘ 3 [Launch into Auto Scaling Group](#) ⓘ

You may want to consider launching these instances into an Auto Scaling Group to help you maintain high availability and help your application stay healthy and cost effective.

**Purchasing option** ⓘ ☐ Request Spot instances

**Network** ⓘ vpc-09bdb16d (10.0.0.0/16) | HDP-VPC [Create new VPC](#)

**Subnet** ⓘ subnet-99e68def (10.0.1.0/24) | 10.0.1.0-us-east-1c [Create new subnet](#)  
250 IP Addresses available

**Auto-assign Public IP** ⓘ Enable

**IAM role** ⓘ None [Create new IAM role](#)

**Shutdown behavior** ⓘ Stop

**Enable termination protection** ⓘ ☐ Protect against accidental termination

**Monitoring** ⓘ ☐ Enable CloudWatch detailed monitoring  
[Additional charges apply.](#)

**EBS-optimized instance** ⓘ ☐ Launch as EBS-optimized instance  
[Additional charges apply.](#)

Figure 5.1

- From Figure 4.3 in Section 4, we click “Launch” to launch the HDP\_AMI base image.
- Choose “m3.xlarge” instance type with 4 vCPUs and 15 GB of RAM.
- Next, you will need to configure your cluster instances.
- Enter “3” in “Number of Instances” field. Then, select HDP\_VPC as the preferred Network to configure networking for your cluster.
- For the subnet field, select the subnet you created in Section 2.1 from the drop down.
- Select “Enable” from the drop down in “Auto-assign Public IP” as highlighted in Figure 5.1.

In the same page, “Configure Instance Details” scroll down to “Advanced Details”.

In “User Data” text field, type the following:

```
sudo mkdir /grid
sudo mkfs -t ext4 /dev/xvdb
sudo mount /dev/xvdb /grid
```

▼ Network interfaces ⓘ

Device	Network Interface	Subnet	Primary IP	Secondary IP addresses
eth0	New network interface ▼	subnet-3745a20i ▼	Auto-assign	Add IP

Add Device

▼ Advanced Details

User data ⓘ ☒ As text ☐ As file ☐ Input is already base64 encoded

```
sudo mkdir /grid
sudo mkfs -t ext4 /dev/xvdb
sudo mount /dev/xvdb /grid
```

Cancel Previous Review and Launch Next: Add Storage

Figure 5.2

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Tag Instance 6. Configure Security Group 7. Review

### Step 6: Configure Security Group

A security group is a set of firewall rules that control the traffic for your instance. On this page, you can add rules to allow specific traffic to reach your instance. For example, if you want to set up a web instance, add rules that allow unrestricted access to the HTTP and HTTPS ports. You can create a new security group or select from an existing one below. [Learn more](#) about Amazon EC2 security groups.

Assign a security group: ☐ Create a new security group ☒ Select an existing security group

Security Group ID	Name	Description
<input type="checkbox"/> sg-45eb953d	default	default VPC security group
<input checked="" type="checkbox"/> sg-04c3bd7c	HDP-AMI-SG	HDP-AMI-SG created 2016-03-28T11:46:08.204-04:00

Inbound rules for sg-04c3bd7c (Selected security groups: sg-04c3bd7c)

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ
All TCP	TCP	0 - 65535	0.0.0.0/0
All traffic	All	All	0.0.0.0/0
SSH	TCP	22	0.0.0.0/0
All UDP	UDP	0 - 65535	0.0.0.0/0

Figure 5.3

Next, you need to configure “Security Group” for the cluster. You had already created a security group, when you prepared the instance in Section 2.2.

You can re-use the same, instead of configuring a new security group, as highlighted in Figure 5.3.

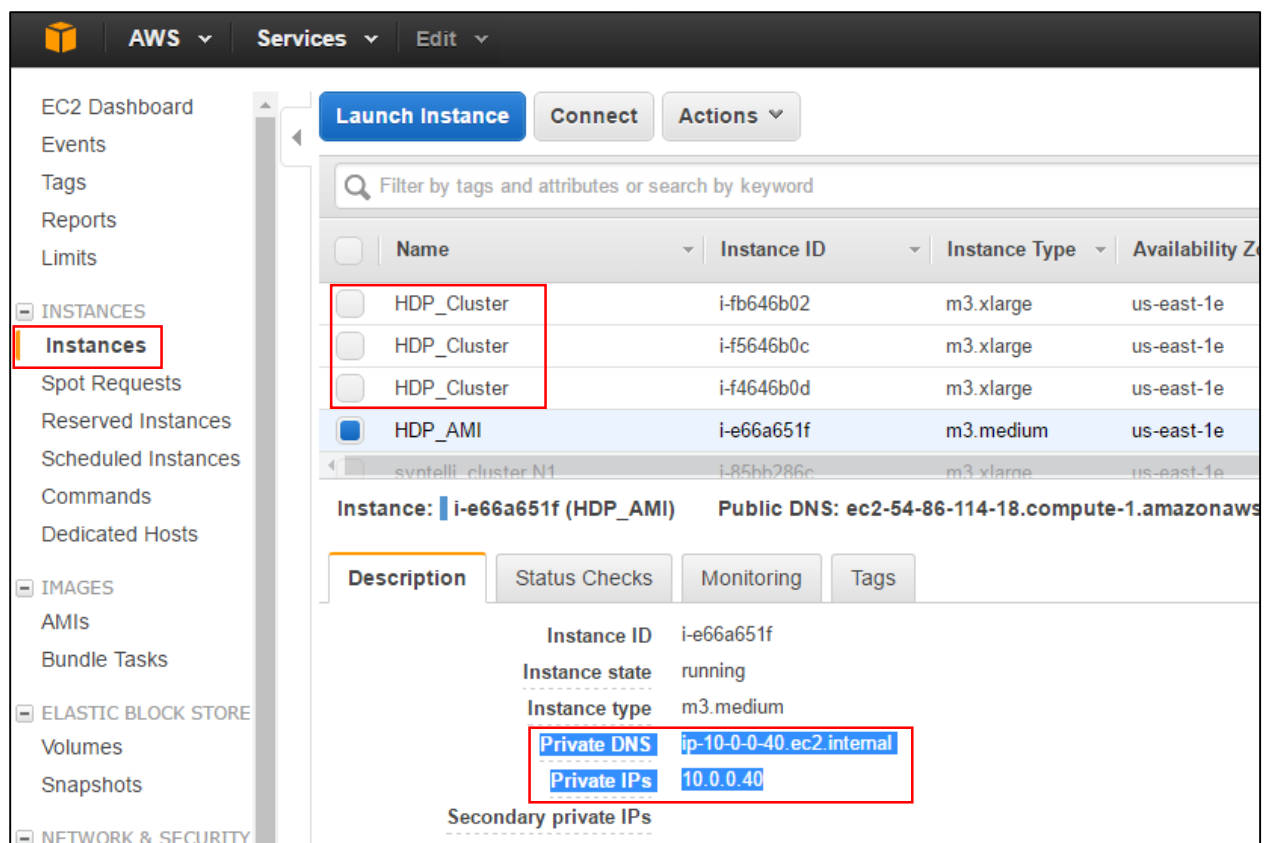


Figure 5.4

NOTE:

- In Figure 5.4, you will notice that HDP\_AMI is an instance of type “m3.medium”. I had created two AMI with different instance type “m3.medium” and “m3.xlarge”.
- I realized that “m3.xlarge” instance type was more effective.

Once, the instances are launched, you can navigate to “Instances” under EC2 Dashboard and you will notice 3 new instances in “**running**” state.

Make a note of the **Private DNS** and **Private IPs** for 3 newly created HDP instances, as highlighted in Figure 5.4.

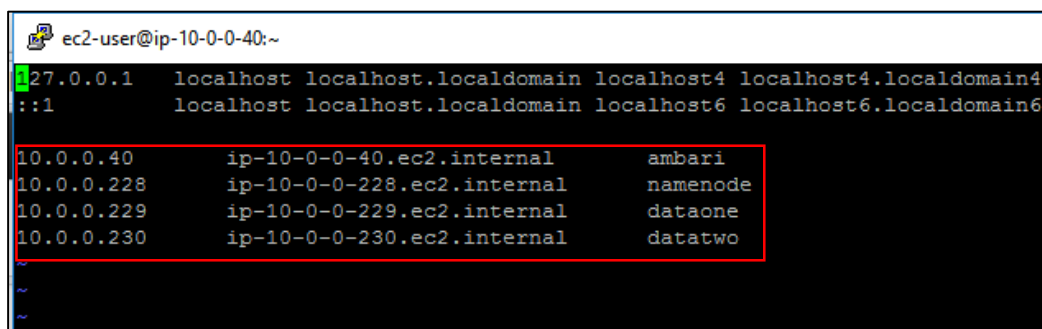


Figure 5.5

SSH as “ec2-user” into HDP\_AMI which will serve as Ambari Server.

Ambari makes Hadoop management simpler by providing a consistent, secure platform for operational control. Ambari provides an intuitive Web UI as well as a robust REST API, which is particularly useful for automating cluster operations.

To learn more about Ambari, visit: <http://hortonworks.com/hadoop/ambari/>

Edit “/etc/hosts” file on all four instances.

```
$ sudo vi /etc/hosts
```

The “**hosts**” file’s contents looks similar to Figure 5.5. You will need to add the Private IP addresses and Private DNS of all four instances:

- a. HDP\_AMI – ambari
- b. HDP\_Cluster – namenode
- c. HDP\_Cluster – dataone
- d. HDP\_Cluster – datatwo

I named the HDP\_Cluster instances namenode, dataone and datatwo based on the ascending numbers of the instance’s Private IP address.

NOTE:

- Your instance IP addresses will differ from the one shown in Figures.
- Make sure, you add the Private IP addresses and Private DNS of all the four instances in “**/etc/hosts**” file for all four instances.

If you try to SSH to one of the data instances from your ambari instance, you should be able login into the instances password-less.

This is possible only, if you have configured the “/etc/hosts” file right.

Also, the SSH Keygen that was created in Section 3.5, helps you log-in to instances on the cluster, not having you to enter password. This is very important for HDP cluster deployment.

To have Ambari Server automatically install Ambari Agents on all your cluster hosts, you must set up password-less SSH connections between the Ambari Server host and all other hosts in the cluster. The Ambari Server host uses SSH public key authentication to remotely access and install the Ambari Agent.

To read more about, password-less SSH on Ambari cluster, visit:

[https://docs.hortonworks.com/HDPDocuments/Ambari-2.2.0.0/bk\\_Installing\\_HDP\\_AMB/content/\\_set\\_up\\_password-less\\_ssh.html](https://docs.hortonworks.com/HDPDocuments/Ambari-2.2.0.0/bk_Installing_HDP_AMB/content/_set_up_password-less_ssh.html)

## 6 AMBARI SERVER INSTALLATION AND SETUP

### 6.1 INSTALL WGET

Install “wget” tool on your Ambari instance. This instance will serve as your Ambari Server.

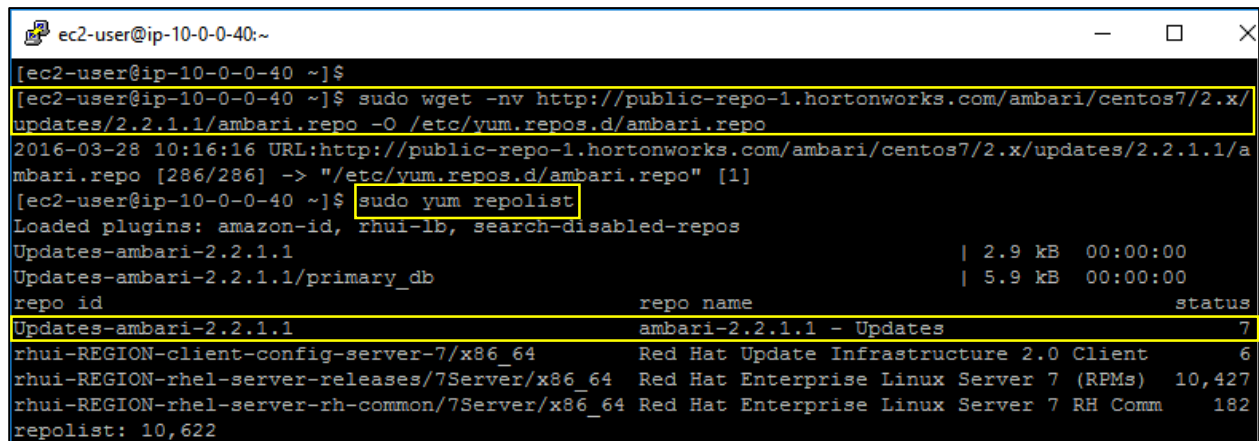
GNU Wget is a free utility for non-interactive download of files from the Web. It supports HTTP, HTTPS, and FTP protocols, as well as retrieval through HTTP proxies.

```
$ sudo yum install wget
```

Next, you need to download the Ambari Repository onto your Ambari Server host.

Make sure, when you copy the above text, it’s a single line execution as highlighted in Figure 6.1.1.

```
$ sudo wget -nv http://public-repo-1.hortonworks.com/ambari/centos7/2.x/updates/2.2.1.1/ambari.repo -O /etc/yum.repos.d/ambari.repo
```



```
ec2-user@ip-10-0-0-40:~$  
[ec2-user@ip-10-0-0-40 ~]$  
[ec2-user@ip-10-0-0-40 ~]$ sudo wget -nv http://public-repo-1.hortonworks.com/ambari/centos7/2.x/updates/2.2.1.1/ambari.repo -O /etc/yum.repos.d/ambari.repo  
2016-03-28 10:16:16 URL:http://public-repo-1.hortonworks.com/ambari/centos7/2.x/updates/2.2.1.1/ambari.repo [286/286] -> "/etc/yum.repos.d/ambari.repo" [1]  
[ec2-user@ip-10-0-0-40 ~]$ sudo yum repolist  
Loaded plugins: amazon-id, rhui-lb, search-disabled-repos  
Updates-ambari-2.2.1.1 | 2.9 kB 00:00:00  
Updates-ambari-2.2.1.1/primary_db | 5.9 kB 00:00:00  
repo id repo name status  
Updates-ambari-2.2.1.1 ambari-2.2.1.1 - Updates 7  
rhui-REGION-client-config-server-7/x86_64 Red Hat Update Infrastructure 2.0 Client 6  
rhui-REGION-rhel-server-releases/7Server/x86_64 Red Hat Enterprise Linux Server 7 (RPMs) 10,427  
rhui-REGION-rhel-server-rh-common/7Server/x86_64 Red Hat Enterprise Linux Server 7 RH Comm 182  
repolist: 10,622
```

Figure 6.1.1

You can verify, if the Ambari Repository has been successfully added to your system repository by issuing the following command.

```
$ sudo yum repolist
```

Now, you need to install the Ambari Server on your “Ambari Server” host. Issue the following command:

```
$ sudo yum install ambari-server
```

When you install the Ambari Server, you will get a token that verifies the installation as shown in Figure 6.1.2.

```
Retrieving key from http://public-repo-1.hortonworks.com/ambari/centos7/RPM-GPG-KEY/RPM-GPG-KEY-J
enkins
Importing GPG key 0x07513CAD:
  Userid      : "Jenkins (HDP Builds) <jenkin@hortonworks.com>"
  Fingerprint: df52 ed4f 7a3a 5882 c099 4c66 b973 3a7a 0751 3cad
  From        : http://public-repo-1.hortonworks.com/ambari/centos7/RPM-GPG-KEY/RPM-GPG-KEY-Jenkins
Is this ok [y/N]: y
```

Figure 6.1.2

## 6.2 SETUP AMBARI SERVER

Before starting the Ambari Server, you must set up the Ambari Server. Setup configures Ambari to talk to the Ambari database, installs the JDK and allows you to customize the user account the Ambari Server daemon will run as. The `ambari-server setup` command manages the setup process. Run the following command on the Ambari server host to start the setup process.

To learn more about ambari server setup, visit:

[http://docs.hortonworks.com/HDPDocuments/Ambari-2.2.1.1/bk\\_Installing\\_HDP\\_AMB/content/\\_set\\_up\\_the\\_ambari\\_server.html](http://docs.hortonworks.com/HDPDocuments/Ambari-2.2.1.1/bk_Installing_HDP_AMB/content/_set_up_the_ambari_server.html)

```
$ sudo ambari-server setup
```

Figure 6.2.1 and 6.2.2 highlight the `ambari-server setup`.

```
ec2-user@ip-10-0-1-12:~$ sudo ambari-server setup
Using python /usr/bin/python
Setup ambari-server
Checking SELinux...
SELinux status is 'disabled'
Customize user account for ambari-server daemon [y/n] (n)? y
Enter user account for ambari-server daemon (root):
Adjusting ambari-server permissions and ownership...
Checking firewall status...
Redirecting to /bin/systemctl status iptables.service

Checking JDK...
[1] Oracle JDK 1.8 + Java Cryptography Extension (JCE) Policy Files 8
[2] Oracle JDK 1.7 + Java Cryptography Extension (JCE) Policy Files 7
[3] Custom JDK
=====
Enter choice (1): 1
To download the Oracle JDK and the Java Cryptography Extension (JCE) Policy Files you must accept
the license terms found at http://www.oracle.com/technetwork/java/javase/terms/license/index.htm
l and not accepting will cancel the Ambari Server setup and you must install the JDK and JCE file
s manually.
Do you accept the Oracle Binary Code License Agreement [y/n] (y)? y
Downloading JDK from http://public-repo-1.hortonworks.com/ARTIFACTS/jdk-8u60-linux-x64.tar.gz to
/var/lib/ambari-server/resources/jdk-8u60-linux-x64.tar.gz
jdk-8u60-linux-x64.tar.gz... 100% (172.8 MB of 172.8 MB)
Successfully downloaded JDK distribution to /var/lib/ambari-server/resources/jdk-8u60-linux-x64.t
ar.gz
Installing JDK to /usr/jdk64/
Successfully installed JDK to /usr/jdk64/
Downloading JCE Policy archive from http://public-repo-1.hortonworks.com/ARTIFACTS/jce_policy-8.z
ip to /var/lib/ambari-server/resources/jce_policy-8.zip
Successfully downloaded JCE Policy archive to /var/lib/ambari-server/resources/jce_policy-8.zip
Installing JCE policy...
Completing setup...
Configuring database...
Enter advanced database configuration [y/n] (n)?
```

Figure 6.2.1

```
Installing JCE policy...
Completing setup...
Configuring database...
Enter advanced database configuration [y/n] (n)? n
Configuring database...
Default properties detected. Using built-in database.
Configuring ambari database...
Checking PostgreSQL...
Running initdb: This may take upto a minute.
Initializing database ... OK

About to start PostgreSQL
Configuring local database...
Connecting to local database...done.
Configuring PostgreSQL...
Restarting PostgreSQL
Extracting system views...
ambari-admin-2.2.1.1.70.jar
.....
Adjusting ambari-server permissions and ownership...
Ambari Server 'setup' completed successfully.
[ec2-user@ip-10-0-1-12 ~]$
```

Figure 6.2.2



```
ec2-user@ip-10-0-0-40:~  
[ec2-user@ip-10-0-0-40 ~]$ sudo ambari-server start  
Using python /usr/bin/python  
Starting ambari-server  
Ambari Server running with administrator privileges.  
Organizing resource files at /var/lib/ambari-server/resources...  
Server PID at: /var/run/ambari-server/ambari-server.pid  
Server out at: /var/log/ambari-server/ambari-server.out  
Server log at: /var/log/ambari-server/ambari-server.log  
Waiting for server start.....  
Ambari Server 'start' completed successfully.  
[ec2-user@ip-10-0-0-40 ~]$
```

Figure 6.2.3

Once you successfully completed, ambari-server setup, you can start the server and proceed to installation process on using the Ambari Web based installation.

```
$ sudo ambari-server start
```

The ambari-server will start on the host, and display few log messages as highlighted in Figure 6.2.3.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public IP
HDP_NN	i-b42b1630	m3.xlarge	us-east-1c	running	2/2 checks ...	None	
<b>HDP_Ambari</b>	<b>i-b52b1631</b>	m3.xlarge	us-east-1c	running	2/2 checks ...	None	
HDP_DN2	i-bb2b163f	m3.xlarge	us-east-1c	running	2/2 checks ...	None	
HDP_DN1	i-cc9db48	m3.xlarge	us-east-1c	running	2/2 checks ...	None	

Instance: i-b52b1631 (HDP_Ambari)		Public IP: 52.207.242.94
Description	Status Checks	Monitoring
Instance ID	i-b52b1631	Public DNS
Instance state	running	Public IP

Figure 6.2.4

## 7 INSTALLING, CONFIGURING, AND DEPLOYING A HDP CLUSTER

---

Use the Ambari Install Wizard running in your browser to install, configure, and deploy your cluster, as follows:

### 7.1 LOG IN TO APACHE AMBARI

After starting the Ambari service, open Ambari Web using a web browser.

Point your browser to `http://<your.ambari.server.external.ip>:8080`, where

`<your.ambari.server.external.ip>` is the name of your ambari server host.

Log in to the Ambari Server using the default user name/password: `admin/admin`. You can change these credentials later. Highlighted in Figure 7.1.1.

For a new cluster, the Ambari install wizard displays a Welcome page from which you launch the Ambari Install wizard.

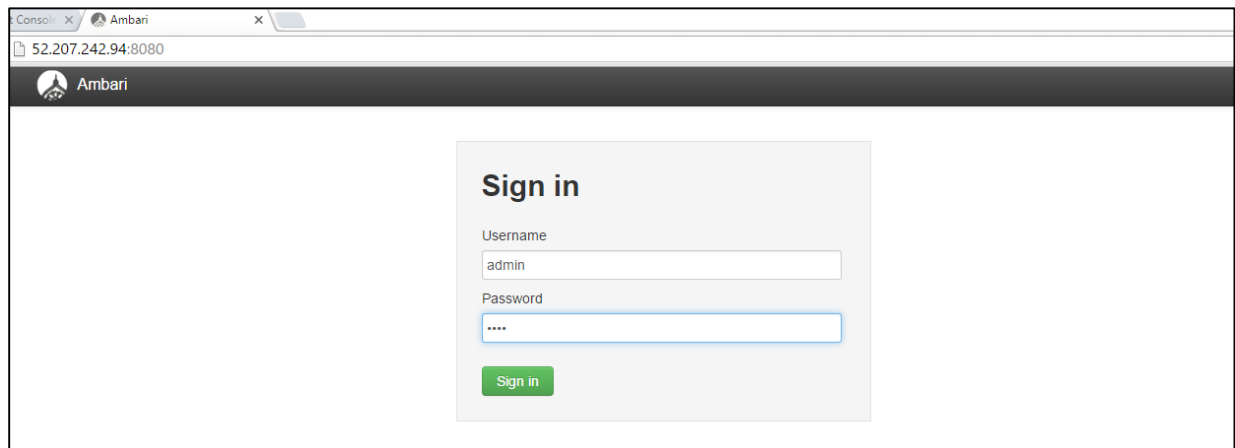


Figure 7.1.1

### 7.2 NAME YOUR CLUSTER

Once, you sign into the Ambari Install Wizard, then you will be presented to a screen shown in Figure 7.2.1.

Click “Launch Install Wizard” to create a new cluster.

You will be navigated to “Get Started” screen, where you need to provide a name for your cluster.

Provide a name for your cluster, and click “Next” as highlighted in Figure 7.2.2.

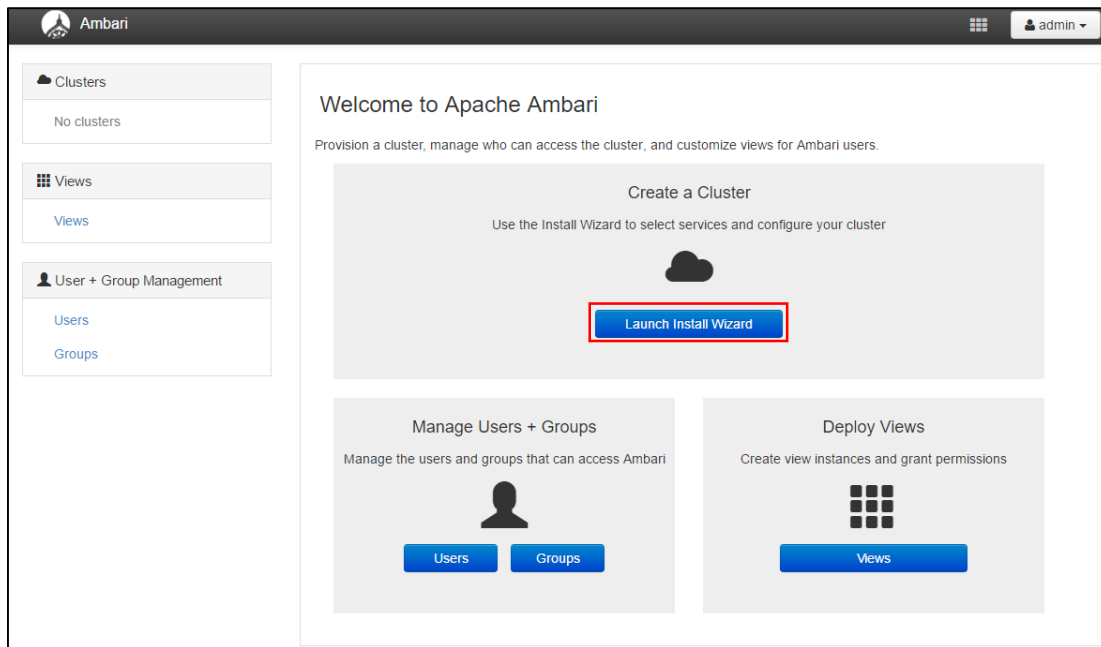


Figure 7.2.1

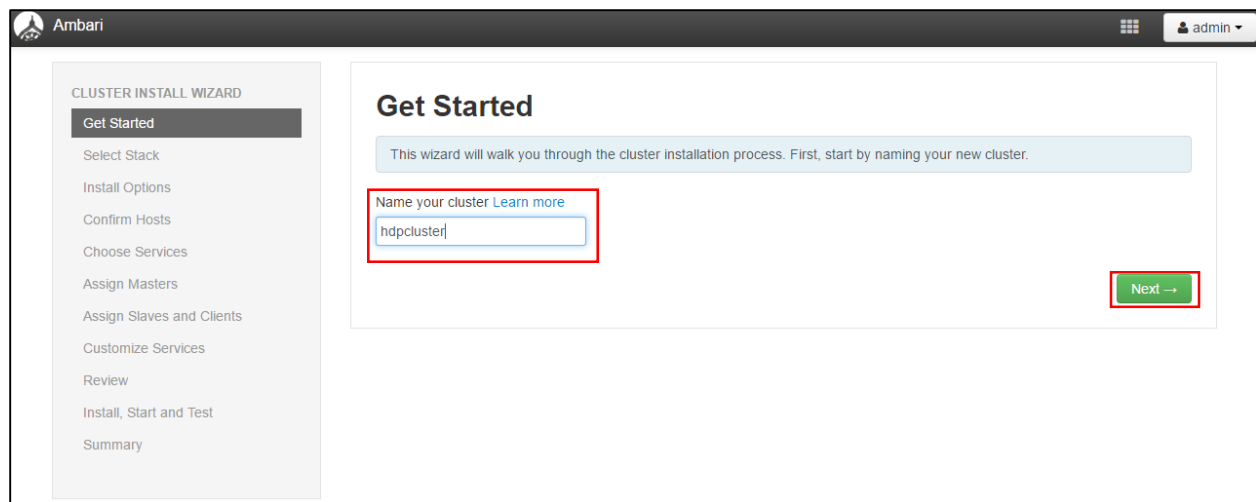


Figure 7.2.2

## 7.3 SELECT STACK

Select HDP stack version, which you would like to install on your cluster.

I selected the latest version to be installed on my AWS cluster, as highlighted in Figure 7.3.1

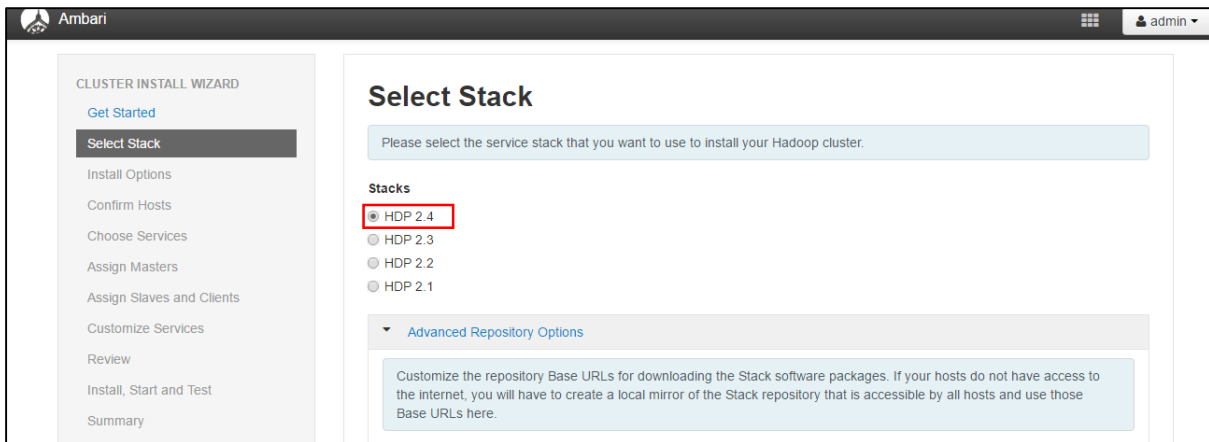


Figure 7.3.1

## 7.4 INSTALL OPTIONS

Select the HDP version for your OS type. I selected redhat7 repositories as highlighted in Figure 7.4.1.

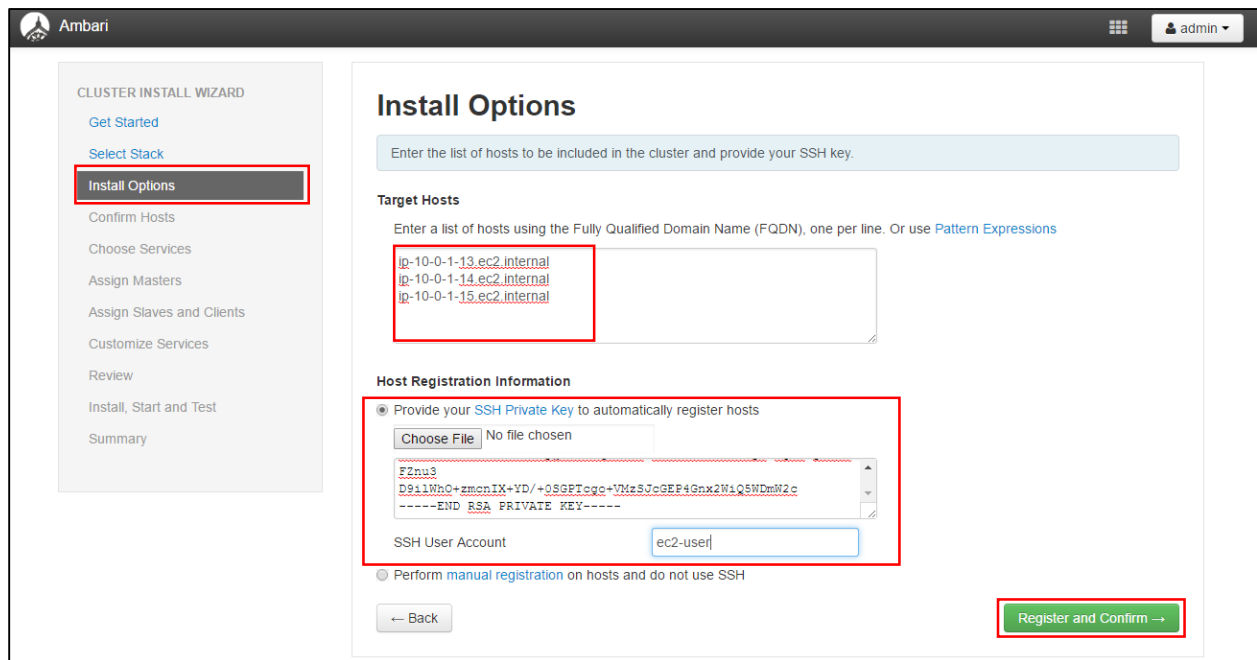
**Important:** When using local mirror repositories, you only need to provide Base URLs for the Operating System you are installing for your Stack. Uncheck all other repositories.

OS	Name	Base URL
<input type="checkbox"/> debian7	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/debian7/2.x/updates/2.4">http://public-repo-1.hortonworks.com/HDP/debian7/2.x/updates/2.4</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/d">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/d</a> ✕
<input type="checkbox"/> redhat6	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.4">http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.4</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/ci">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/ci</a> ✕
<input checked="" type="checkbox"/> redhat7	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/centos7/2.x/updates/2.4">http://public-repo-1.hortonworks.com/HDP/centos7/2.x/updates/2.4</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/ci">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/ci</a> ✕
<input type="checkbox"/> suse11	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/suse11sp3/2.x/updates/">http://public-repo-1.hortonworks.com/HDP/suse11sp3/2.x/updates/</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/si">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/si</a> ✕
<input type="checkbox"/> ubuntu12	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2">http://public-repo-1.hortonworks.com/HDP/ubuntu12/2.x/updates/2</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/u">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/u</a> ✕
<input type="checkbox"/> ubuntu14	HDP-2.4	<a href="http://public-repo-1.hortonworks.com/HDP/ubuntu14/2.x/updates/2">http://public-repo-1.hortonworks.com/HDP/ubuntu14/2.x/updates/2</a> ✕
	HDP-UTILS-1.1.0.20	<a href="http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/u">http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/u</a> ✕

☐ Skip Repository Base URL validation (Advanced) ?

← Back Next →

Figure 7.4.1

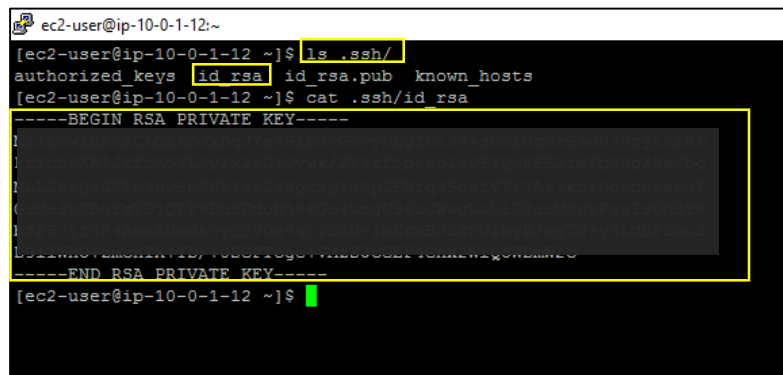


The screenshot shows the Ambari web interface during the cluster installation wizard. The left sidebar contains the 'CLUSTER INSTALL WIZARD' menu with options: Get Started, Select Stack, **Install Options** (highlighted with a red box), Confirm Hosts, Choose Services, Assign Masters, Assign Slaves and Clients, Customize Services, Review, Install, Start and Test, and Summary. The main content area is titled 'Install Options' and includes a text box for 'Enter the list of hosts to be included in the cluster and provide your SSH key.' Below this is the 'Target Hosts' section with a text area containing three hostnames: ip-10-0-1-13.ec2.internal, ip-10-0-1-14.ec2.internal, and ip-10-0-1-15.ec2.internal (highlighted with a red box). The 'Host Registration Information' section has two radio buttons: 'Provide your SSH Private Key to automatically register hosts' (selected) and 'Perform manual registration on hosts and do not use SSH'. Under the selected option, there is a 'Choose File' button (disabled), a text area containing an SSH private key (highlighted with a red box), and an 'SSH User Account' field with 'ec2-user' entered. At the bottom right is a green 'Register and Confirm -->' button (highlighted with a red box).

Figure 7.4.2

Next, enter the Private DNS of instances of namenode, dataone and datatwo. One instance will be configured as Name Node and the other two will be configured as Data Nodes.

Copy the private key from “id\_rsa” file on you Ambari Server host. This file is located in “.ssh” directory.



```
ec2-user@ip-10-0-1-12:~$ ls .ssh/
authorized_keys  id_rsa  id_rsa.pub  known_hosts
ec2-user@ip-10-0-1-12 ~$ cat .ssh/id_rsa
-----BEGIN RSA PRIVATE KEY-----
[REDACTED]
-----END RSA PRIVATE KEY-----
ec2-user@ip-10-0-1-12 ~$
```

The terminal screenshot shows the user 'ec2-user' on host 'ip-10-0-1-12'. They run 'ls .ssh/' showing files 'authorized\_keys', 'id\_rsa', 'id\_rsa.pub', and 'known\_hosts'. Then they run 'cat .ssh/id\_rsa', which displays the contents of the private key file, from '-----BEGIN RSA PRIVATE KEY-----' to '-----END RSA PRIVATE KEY-----'. The key content is redacted with a black box (highlighted with a yellow box in the original image).

Figure 7.4.3

Copy it from “---BEGIN RSA PRIVATE KEY--” to “----END RSA PRIVATE KEY---” as highlighted in Figure 7.4.3.

Finally, click “Register and Confirm”.

## 7.5 CONFIRM HOSTS

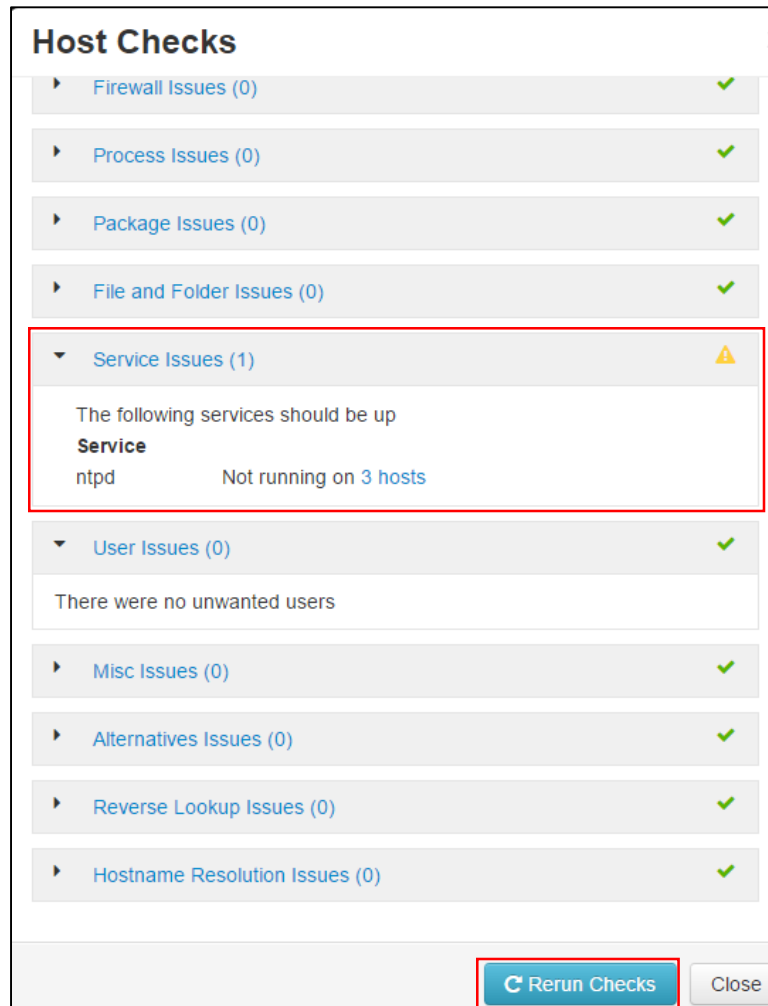


Figure 7.5.1

If the hosts have any issue, Ambari installation wizard will highlight it to the user as shown in Figure 7.5.1

In my case, I had to enable “ntpd” (NTP daemon). I ran the following commands shown in Figure 7.5.2.

```
// Check Hosts
The following services should be up
Service
ntpd Not running on 3 hosts
.....

$ sudo systemctl enable ntpd
.....
$ sudo systemctl start ntpd
.....
$ sudo systemctl is-enabled ntpd
.....
```

Figure 7.5.2

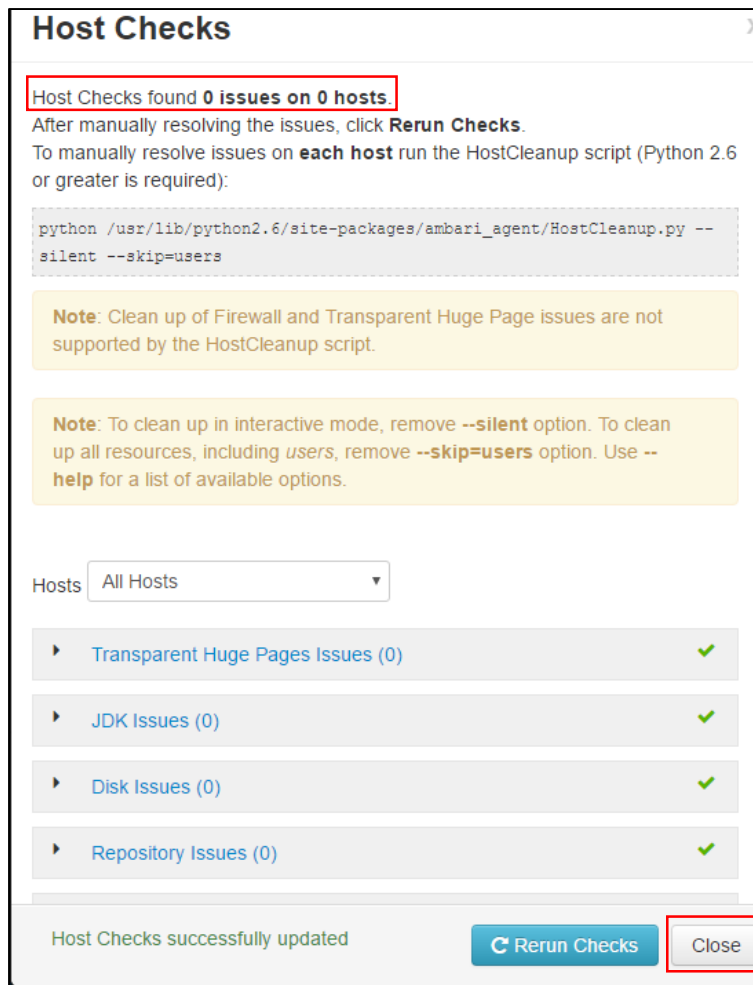


Figure 7.5.3

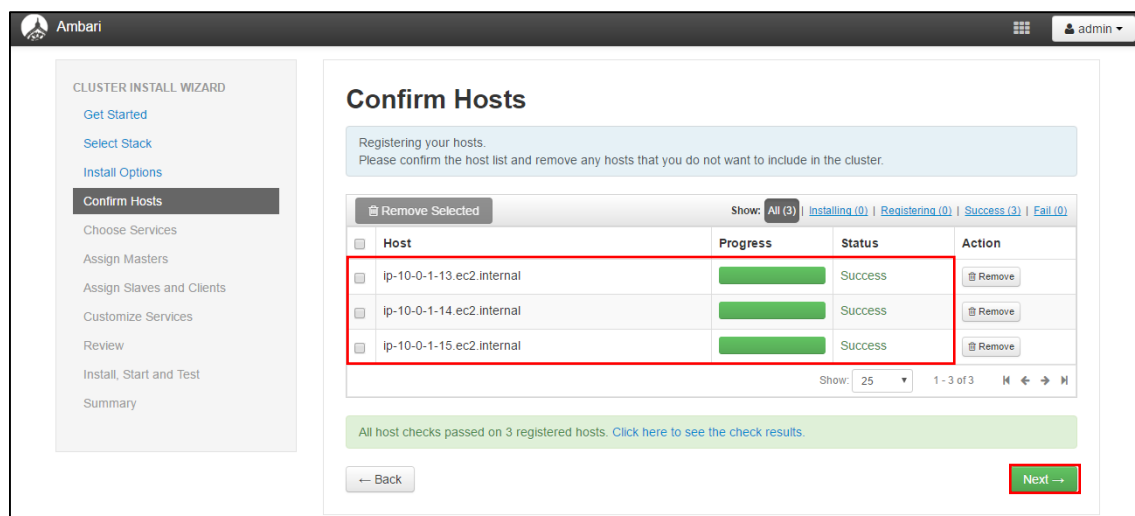
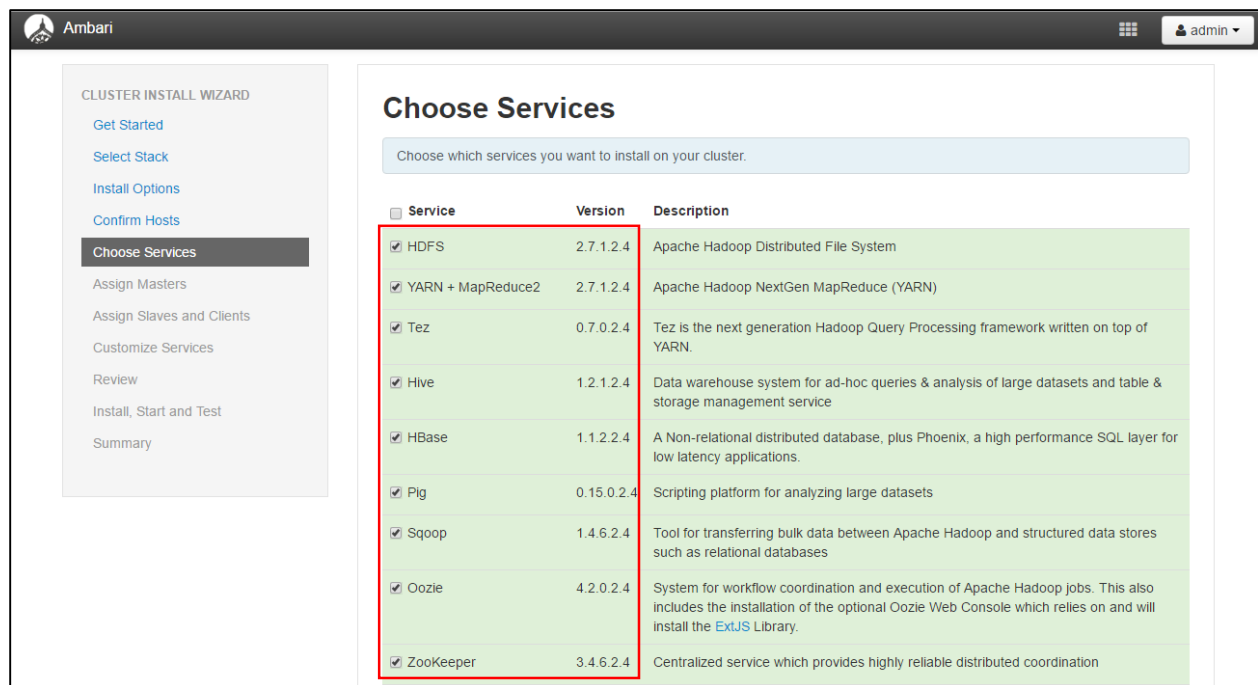


Figure 7.5.4

If the Ambari Installation Wizard, does not find any issues with the hosts, it report zero errors as highlighted in Figure 7.5.3.

Also successful registration of the hosts by the Ambari Server is checked “green” as highlighted in Figure 7.5.4.

## 7.6 CHOOSE SERVICES



The screenshot shows the Ambari web interface. On the left is a sidebar with the 'CLUSTER INSTALL WIZARD' steps: Get Started, Select Stack, Install Options, Confirm Hosts, **Choose Services** (highlighted), Assign Masters, Assign Slaves and Clients, Customize Services, Review, Install, Start and Test, and Summary. The main content area is titled 'Choose Services' and contains a text box: 'Choose which services you want to install on your cluster.' Below this is a table with columns: Service, Version, and Description. A red box highlights the first eight rows of the table.

<input type="checkbox"/> Service	Version	Description
<input checked="" type="checkbox"/> HDFS	2.7.1.2.4	Apache Hadoop Distributed File System
<input checked="" type="checkbox"/> YARN + MapReduce2	2.7.1.2.4	Apache Hadoop NextGen MapReduce (YARN)
<input checked="" type="checkbox"/> Tez	0.7.0.2.4	Tez is the next generation Hadoop Query Processing framework written on top of YARN.
<input checked="" type="checkbox"/> Hive	1.2.1.2.4	Data warehouse system for ad-hoc queries & analysis of large datasets and table & storage management service
<input checked="" type="checkbox"/> HBase	1.1.2.2.4	A Non-relational distributed database, plus Phoenix, a high performance SQL layer for low latency applications.
<input checked="" type="checkbox"/> Pig	0.15.0.2.4	Scripting platform for analyzing large datasets
<input checked="" type="checkbox"/> Sqoop	1.4.6.2.4	Tool for transferring bulk data between Apache Hadoop and structured data stores such as relational databases
<input checked="" type="checkbox"/> Oozie	4.2.0.2.4	System for workflow coordination and execution of Apache Hadoop jobs. This also includes the installation of the optional Oozie Web Console which relies on and will install the <a href="#">ExtJS Library</a> .
<input checked="" type="checkbox"/> ZooKeeper	3.4.6.2.4	Centralized service which provides highly reliable distributed coordination

Figure 7.6.1



<input type="checkbox"/> Storm	0.10.0.2.4	Apache Hadoop Stream processing framework
<input checked="" type="checkbox"/> Flume	1.5.2.2.4	A distributed service for collecting, aggregating, and moving large amounts of streaming data into HDFS
<input type="checkbox"/> Accumulo	1.7.0.2.4	Robust, scalable, high performance distributed key/value store.
<input checked="" type="checkbox"/> Ambari Metrics	0.1.0	A system for metrics collection that provides storage and retrieval capability for metrics collected from the cluster
<input type="checkbox"/> Atlas	0.5.0.2.4	Atlas Metadata and Governance platform
<input checked="" type="checkbox"/> Kafka	0.9.0.2.4	A high-throughput distributed messaging system
<input type="checkbox"/> Knox	0.6.0.2.4	Provides a single point of authentication and access for Apache Hadoop services in a cluster
<input checked="" type="checkbox"/> Mahout	0.9.0.2.4	Project of the Apache Software Foundation to produce free implementations of distributed or otherwise scalable machine learning algorithms focused primarily in the areas of collaborative filtering, clustering and classification
<input type="checkbox"/> Slider	0.80.0.2.4	A framework for deploying, managing and monitoring existing distributed applications on YARN.
<input type="checkbox"/> SmartSense	1.2.1.0-70	SmartSense - Hortonworks SmartSense Tool (HST) helps quickly gather configuration, metrics, logs from common HDP services that aids to quickly troubleshoot support cases and receive cluster-specific recommendations.
<input checked="" type="checkbox"/> Spark	1.6.0.2.4	Apache Spark is a fast and general engine for large-scale data processing.

← Back
Next →

Figure 7.6.2

I chose services and tools for my cluster. You can choose, the tools you would like to test on your cluster.

Click “Next”.

## 7.7 ASSIGN MASTERS

CLUSTER INSTALL WIZARD

- Get Started
- Select Stack
- Install Options
- Confirm Hosts
- Choose Services
- Assign Masters**
- Assign Slaves and Clients
- Customize Services
- Review
- Install, Start and Test
- Summary

### Assign Masters

Assign master components to hosts you want to run them on.  
\* HiveServer2 and WebHCat Server will be hosted on the same host.

SNameNode:

NameNode:

History Server:

App Timeline Server:

ResourceManager:

HiveServer2:

Hive Metastore:

WebHCat Server:

HBase Master:

Oozie Server:

ZooKeeper Server:

ZooKeeper Server:

ip-10-0-1-13.ec2.internal (14.3 GB, 4 cores)

NameNode, HiveServer2, Hive Metastore, WebHCat Server, ZooKeeper Server

ip-10-0-1-14.ec2.internal (14.3 GB, 4 cores)

SNameNode, History Server, ResourceManager, HBase Master, Oozie Server, ZooKeeper Server

ip-10-0-1-15.ec2.internal (14.3 GB, 4 cores)

App Timeline Server, ZooKeeper Server, Kafka Broker, Metrics Collector, Spark History Server

Figure 7.7.1

Assign masters, as to which instance will host, which service.

I spread my Hadoop services as shown in Figure 7.7.1. You can plan your services on the hosts in the same manner.

## 7.8 ASSIGN SLAVES AND CLIENTS

CLUSTER INSTALL WIZARD

Get Started

Select Stack

Install Options

Confirm Hosts

Choose Services

Assign Masters

**Assign Slaves and Clients**

Customize Services

Review

Install, Start and Test

Summary

Assign Slaves and Clients

Assign slave and client components to hosts you want to run them on.

Hosts that are assigned master components are shown with \*.

"Client" will install HDFS Client, MapReduce2 Client, YARN Client, Tez Client, HCat Client, Hive Client, HBase Client, Pig, Sqoop, Oozie Client, ZooKeeper Client, Mahout and Spark Client.

Host	all	none	all	none	all	none	all	none	all	none	all
ip-10-0-1-13.ec2.internal*	<input type="checkbox"/> DataNode	<input checked="" type="checkbox"/> NFSGateway	<input type="checkbox"/> NodeManager	<input checked="" type="checkbox"/> RegionServer	<input checked="" type="checkbox"/> Phoenix Query Server	<input checked="" type="checkbox"/>					
ip-10-0-1-14.ec2.internal*	<input checked="" type="checkbox"/> DataNode	<input type="checkbox"/> NFSGateway	<input checked="" type="checkbox"/> NodeManager	<input type="checkbox"/> RegionServer	<input checked="" type="checkbox"/> Phoenix Query Server	<input type="checkbox"/>					
ip-10-0-1-15.ec2.internal*	<input checked="" type="checkbox"/> DataNode	<input type="checkbox"/> NFSGateway	<input checked="" type="checkbox"/> NodeManager	<input checked="" type="checkbox"/> RegionServer	<input type="checkbox"/> Phoenix Query Server	<input checked="" type="checkbox"/>					

Show: 25 1 - 3 of 3

← Back

Next →

Figure 7.8.1

CLUSTER INSTALL WIZARD

Get Started

Select Stack

Install Options

Confirm Hosts

Choose Services

Assign Masters

**Assign Slaves and Clients**

Customize Services

Review

Install, Start and Test

Summary

Assign Slaves and Clients

Assign slave and client components to hosts you want to run them on.

Hosts that are assigned master components are shown with \*.

"Client" will install HDFS Client, MapReduce2 Client, YARN Client, Tez Client, HCat Client, Hive Client, HBase Client, Pig, Sqoop, Oozie Client, ZooKeeper Client, Mahout and Spark Client.

	none	all	none	all	none	all	none	all	none	all	none
NFSGateway	<input type="checkbox"/> NodeManager	<input checked="" type="checkbox"/> RegionServer	<input checked="" type="checkbox"/> Phoenix Query Server	<input checked="" type="checkbox"/> Flume	<input type="checkbox"/> Spark Thrift Server	<input checked="" type="checkbox"/> Client					
NFSGateway	<input checked="" type="checkbox"/> NodeManager	<input type="checkbox"/> RegionServer	<input checked="" type="checkbox"/> Phoenix Query Server	<input type="checkbox"/> Flume	<input checked="" type="checkbox"/> Spark Thrift Server	<input checked="" type="checkbox"/> Client					
NFSGateway	<input checked="" type="checkbox"/> NodeManager	<input checked="" type="checkbox"/> RegionServer	<input type="checkbox"/> Phoenix Query Server	<input checked="" type="checkbox"/> Flume	<input checked="" type="checkbox"/> Spark Thrift Server	<input checked="" type="checkbox"/> Client					

Show: 25 1 - 3 of 3

← Back

Next →

Figure 7.8.2

Decide, the data node instances on your cluster and accordingly install the services and clients required.

## 7.9 CUSTOMIZE SERVICES

**Customize Services**

We have come up with recommended configurations for the services you selected. Customize them as you see fit.

HDFS MapReduce2 YARN Tez **Hive 1** HBase Pig Sqoop **Oozie 1** ZooKeeper Flume

Amba Metrics Kafka Mahout Spark Misc

Group: Default (3) Manage Config Groups Filter...

Settings Advanced

**Advanced core-site**

hadoop.proxyuser.hcat.hosts \*

hadoop.proxyuser.hive.hosts \*

hadoop.proxyuser.oozie.hosts |

Figure 7.9.1

Figure 7.9.1 shows services highlighted, I had to attend and resolve. You might come across similar issue.

**Customize Services**

We have come up with recommended configurations for the services you selected. Customize them as you see fit.

HDFS MapReduce2 YARN Tez **Hive** HBase Pig Sqoop Oozie ZooKeeper Flume Amba Metrics

Kafka Mahout Spark Misc

Group: Default (3) Manage Config Groups Filter...

Settings Advanced

**Hive Metastore**

Hive Metastore host: ip-10-0-1-13.ec2.internal

**Hive Database**

☒ New MySQL Database

☐ Existing MySQL Database

☐ Existing PostgreSQL Database

☐ Existing Oracle Database

☐ Existing SQL Anywhere Database

Database Host: ip-10-0-1-13.ec2.internal

Database Name: hive

Database Username: hive

Database Password: \*\*\*\*\*

Figure 7.9.2

The Hive Database needed configuration. I created a new MySQL Database for Hive metastore.

Provide the username and password. Make sure, you document or note down the username and password.

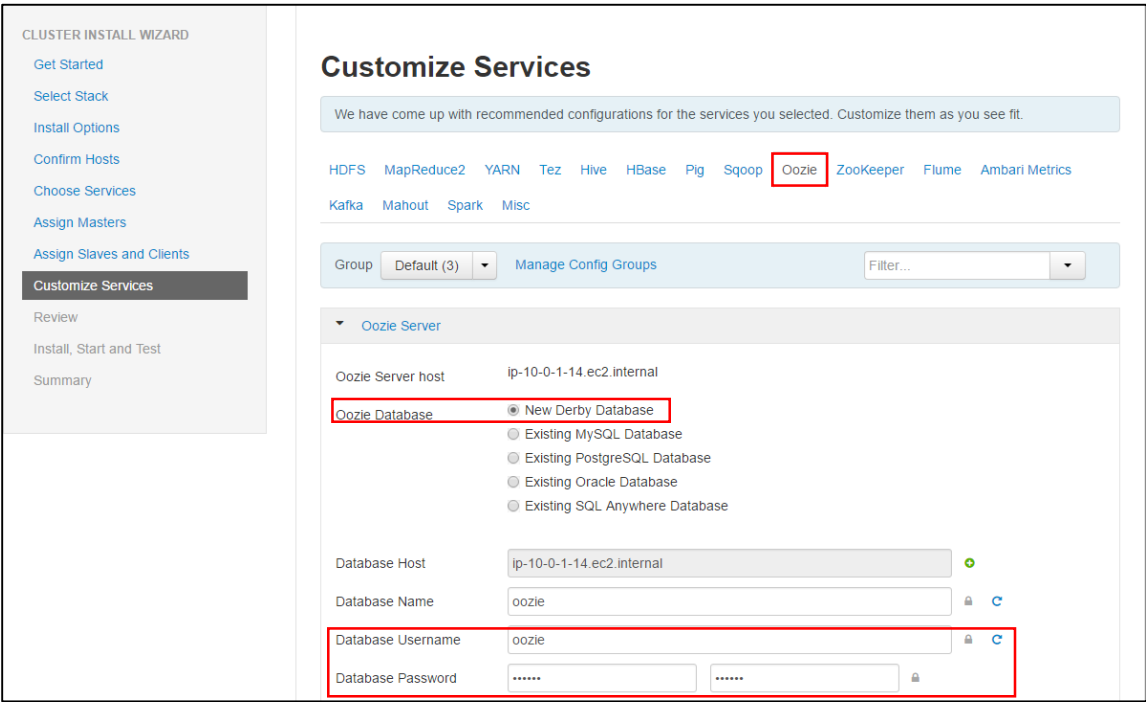


Figure 7.9.3

The Oozie Database needed configuration. I created a new Derby Database for Oozie DB.

Provide the username and password. Make sure, you document or note down the username and password.

NOTE:

- Derby DB is not recommended for production use.

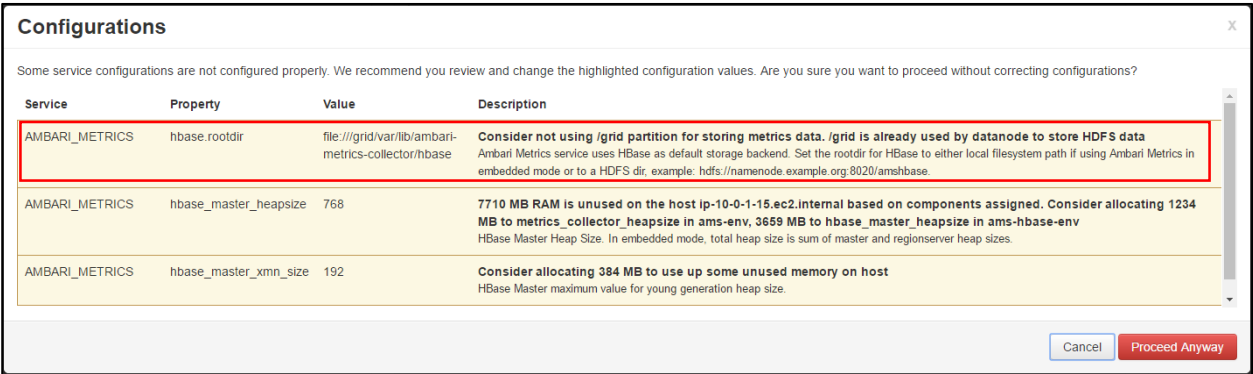


Figure 7.9.4

When you resolve all warnings and click “Proceed”, you might get few more warning with respect to Ambari Metrics memory allocation as shown in Figure 7.9.4.

Ambari Installation Wizard will suggest you to assign Ambari Metrics root directory. You can set it by locating, where the HBase root directory is located as highlighted in Figure 7.9.5

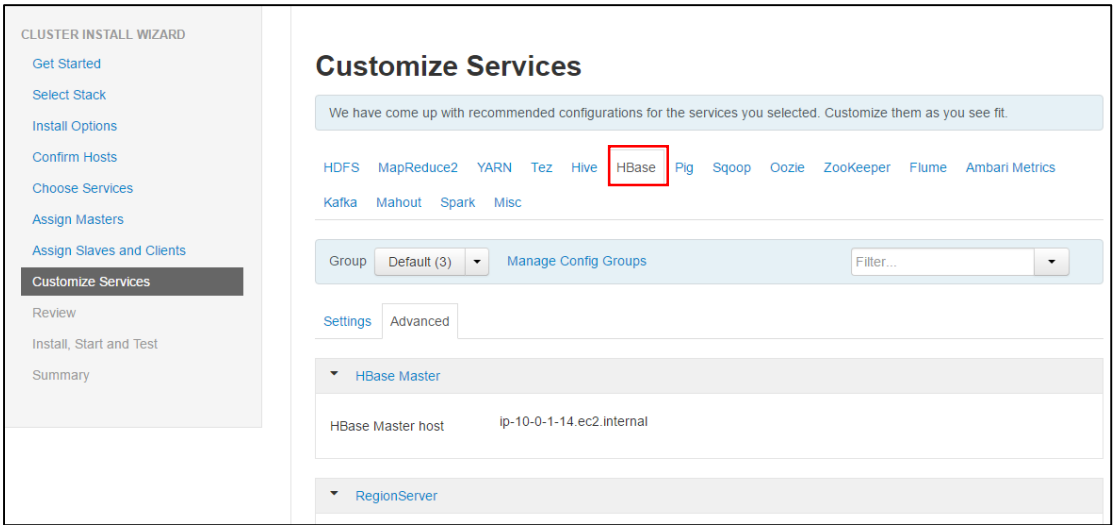


Figure 7.9.5

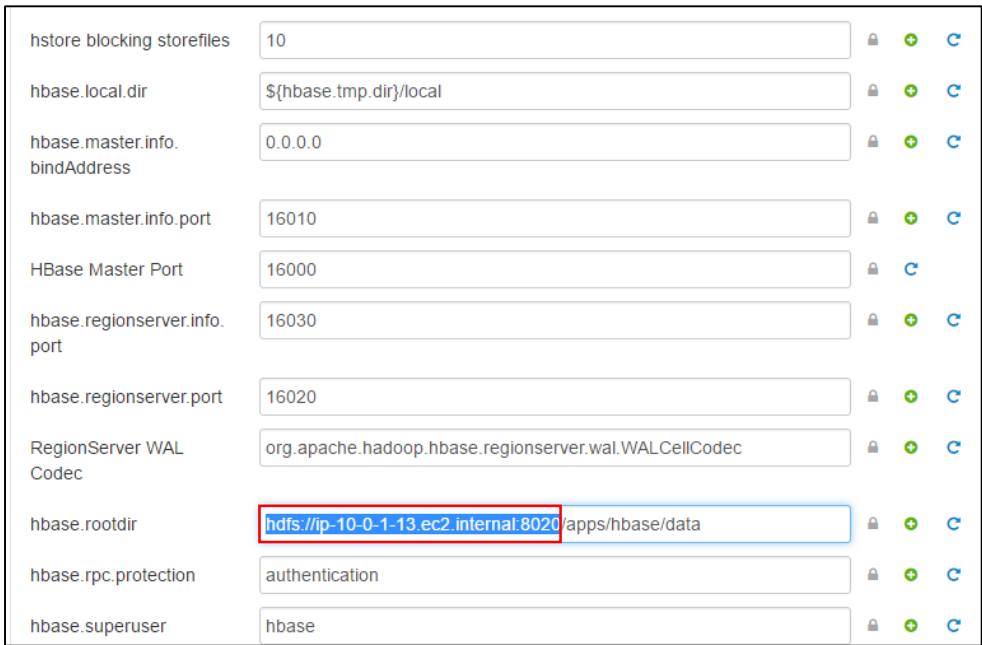


Figure 7.9.6

You can copy the path of your Hbase root directory as highlighted in Figure 7.9.6.

The screenshot shows the 'Customize Services' interface in Ambari. On the left is a sidebar with navigation links: Get Started, Select Stack, Install Options, Confirm Hosts, Choose Services, Assign Masters, Assign Slaves and Clients, **Customize Services**, Review, Install, Start and Test, and Summary. The main content area is titled 'Customize Services' and includes a message: 'We have come up with recommended configurations for the services you selected. Customize them as you see fit.' Below this are tabs for various services: HDFS, MapReduce2, YARN, Tez, Hive, HBase, Pig, Sqoop, Oozie, ZooKeeper, Flume, and Ambari Metrics. A 'Group' dropdown is set to 'Default (3)' with a 'Manage Config Groups' link and a 'Filter...' input. Under the 'General' tab, the 'Metrics Service operation mode' is set to 'embedded', which is highlighted with a red rectangle. To the right of the text field are green and blue circular icons.

Figure 7.9.7

This screenshot shows the same 'Customize Services' page, but the 'Metrics Service operation mode' has been changed to 'distributed', which is highlighted with a red rectangle. Above the configuration fields, a yellow banner states: 'There are 6 configuration changes in 1 service Show Details'. The sidebar and service tabs remain the same. Below the 'Metrics Service operation mode' field, there are four more configuration fields, each with a green and blue icon to its right: 'Metrics Collector log dir' (path: /var/log/ambari-metrics-collector), 'Metrics Collector pid dir' (path: /var/run/ambari-metrics-collector), 'Metrics Monitor log dir' (path: /var/log/ambari-metrics-monitor), and 'Metrics Monitor pid dir' (path: /var/run/ambari-metrics-monitor).

Figure 7.9.8

Change the value for the field “Metrics Service operation mode” from “embedded” to “distributed” highlighted in Figure 7.9.7 to Figure 7.9.8.

Also set “hbase.cluster.distributed” field to “true” as highlighted in Figure 7.9.9.

You will notice ambari metrics root has similar directory path as Hbase root directory, highlighted in Figure 7.9.10.

Advanced ams-hbase-site

dfs.client.read.shortcircuit

true

+

C

hbase.client.scanner.caching

10000

+

C

hbase.client.scanner.timeout.period

900000

+

C

hbase.cluster.distributed

true

+

C

Figure 7.9.9

hbase.rootdir

hdfs://ip-10-0-1-13.ec2.internal:8020/user/ams/hbase

+

C

hbase.snapshot.enabled

false

+

C

hbase.tmp.dir

/var/lib/ambari-metrics-collector/hbase-tmp

+

C

Figure 7.9.10

Set the “hbase\_regionserver\_heapsize” and “regionserver\_xmn\_size” to the recommended value

Advanced ams-hbase-env

hbase\_classpath\_additional

+

hbase\_log\_dir

/var/log/ambari-metrics-collector

+

C

hbase\_master\_heapsize

512

MB

+

C

hbase\_master\_maxperm\_size

128

MB

+

C

hbase\_master\_xmn\_size

102

MB

+

C

hbase\_pid\_dir

/var/run/ambari-metrics-collector/

+

C

hbase\_regionserver\_heapsize

3864

MB

+

C

hbase\_regionserver\_xmn\_ratio

0.2

+

C

max\_open\_files\_limit

32768

+

C

regionserver\_xmn\_size

512

MB

+

C

Figure 7.9.11

Advanced ams-env

metrics\_collector\_heapsize 1268 MB

ams-env template

```

{{ams_hbase_normalizer_enabled}}

# HBase compaction policy enabled
export AMS_HBASE_FIFO_COMPACTION_ENABLED=
{{ams_hbase_fifo_compaction_enabled}}

```

Figure 7.9.12

Configurations

Some service configurations are not configured properly. We recommend you review and change the highlighted configuration values. Are you sure you want to proceed without correcting configurations?

Service	Property	Value	Description
AMBARI_METRICS	hbase_regionserver_heapsize	3864	7210 MB RAM is unused on the host ip-10-0-1-15.ec2.internal based on components assigned. Consider allocating 1890 MB to metrics_collector_heapsize in ams-env, 6355 MB to hbase_regionserver_heapsize in ams-hbase-env. HBase RegionServer Heap Size. In embedded mode, total heap size is sum of master and regionserver heap sizes.
AMBARI_METRICS	regionserver_xmn_size	512	Consider allocating 768 MB to use up some unused memory on host. HBase RegionServer maximum value for young generation heap size.

Cancel Proceed Anyway

Figure 7.9.13

I clicked “Proceed Anyway” after configuring “metrics\_collector\_heapsize” to “1268” highlighted in Figure 7.9.12.

Warnings may still appear. You can look for help by looking at Hortonworks Community forum.

Visit: <http://hortonworks.com/community/>

Warning

Derby is not recommended for production use. With Derby, Oozie Server HA and concurrent connection support will not be available.

Cancel Proceed Anyway

Figure 7.9.14

I mentioned earlier, stating Derby DB is not recommended for production use. When you install, you will also be provided with a similar warning. Click “Proceed Anyway”.



## 7.10 REVIEW

### Review

Please review the configuration before installation

**Admin Name :** admin  
**Cluster Name :** hdpcluster  
**Total Hosts :** 3 (3 new)  
**Repositories:**

- debian7 (HDP-2.4):  
<http://public-repo-1.hortonworks.com/HDP/debian7/2.x/updates/2.4.0.0>
- debian7 (HDP-UTILS-1.1.0.20):  
<http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/debian6>
- redhat6 (HDP-2.4):  
<http://public-repo-1.hortonworks.com/HDP/centos6/2.x/updates/2.4.0.0>
- redhat6 (HDP-UTILS-1.1.0.20):  
<http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/centos6>
- redhat7 (HDP-2.4):  
<http://public-repo-1.hortonworks.com/HDP/centos7/2.x/updates/2.4.0.0>
- redhat7 (HDP-UTILS-1.1.0.20):  
<http://public-repo-1.hortonworks.com/HDP-UTILS-1.1.0.20/repos/centos7>
- suse11 (HDP-2.4):  
<http://public-repo-1.hortonworks.com/HDP/suse11sp3/2.x/updates/2.4.0.0>

[← Back](#) [Print](#) [Deploy →](#)

Figure 7.10.1

Review the services and clients, that will be launched on your cluster on AWS.

## 7.11 INSTALL, START AND TEST

Installation of Hadoop services and Clients is fairly easy, if you have right libraries and binaries.

In my installation process, I had faced the problem of my hosts missing “snappy-devel” libraries.

When, I did my research as to why my installation failed, I realized that HDP repo-list did not install “snappy-devel” library but just “snappy” library. Refer Figure 7.11.1.



Figure 7.11.1

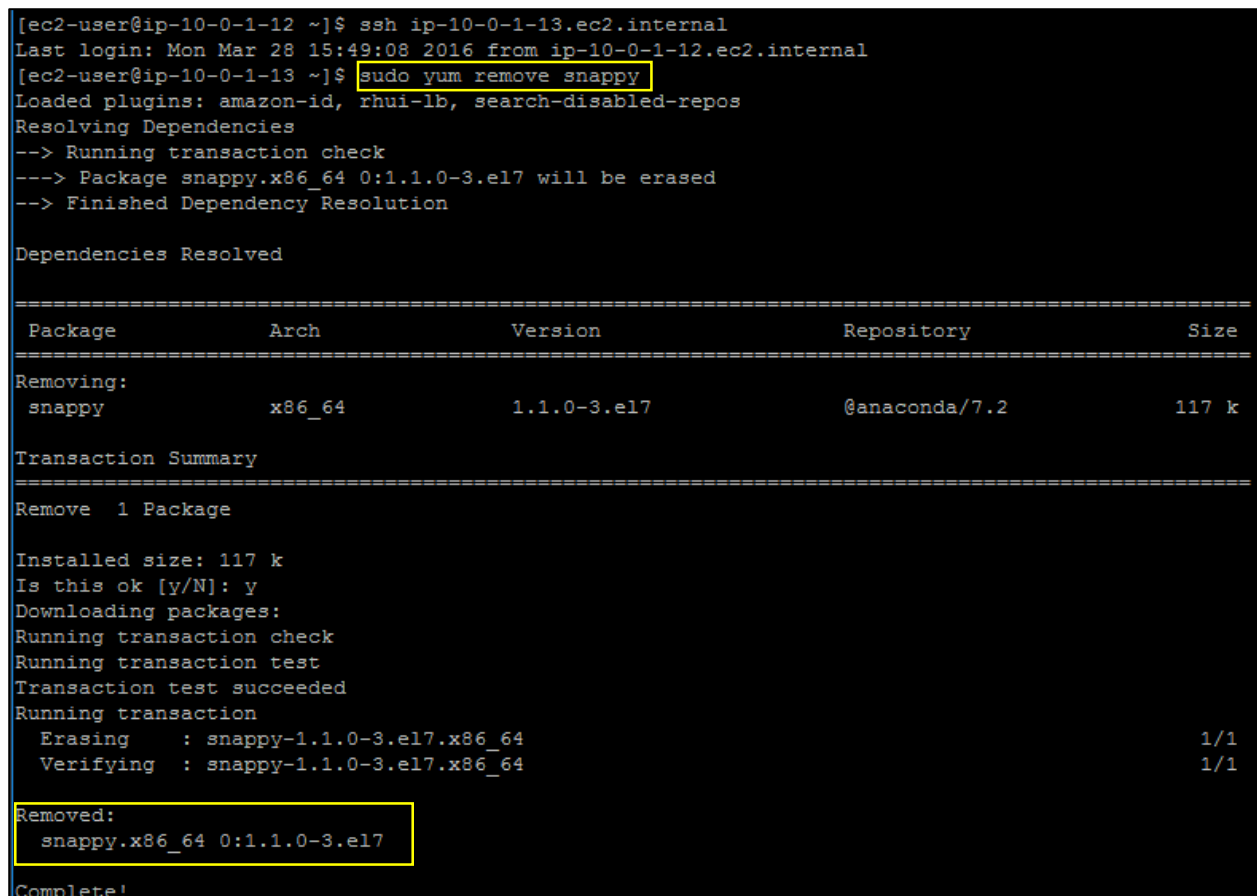


Figure 7.11.2

```
[ec2-user@ip-10-0-1-13 ~]$ sudo yum install snappy-devel
Loaded plugins: amazon-id, rhui-lb, search-disabled-repos
Resolving Dependencies
--> Running transaction check
---> Package snappy-devel.x86_64 0:1.0.5-1.el6 will be installed
--> Processing Dependency: snappy(x86-64) = 1.0.5-1.el6 for package: snappy-devel-1.0.5-1.el6.x86_64
--> Running transaction check
---> Package snappy.x86_64 0:1.0.5-1.el6 will be installed
--> Finished Dependency Resolution

Dependencies Resolved

=====
Package                Arch          Version           Repository        Size
=====
Installing:
snappy-devel            x86_64        1.0.5-1.el6       HDP-UTILS-1.1.0.20 12 k
Installing for dependencies:
snappy                  x86_64        1.0.5-1.el6       HDP-UTILS-1.1.0.20 34 k
=====

Transaction Summary
=====
Install 1 Package (+1 Dependent package)

Total download size: 45 k
Installed size: 112 k
Is this ok [y/d/N]: y
Downloading packages:
(1/2): snappy-devel-1.0.5-1.el6.x86_64.rpm | 12 kB 00:00:00
(2/2): snappy-1.0.5-1.el6.x86_64.rpm      | 34 kB 00:00:00
-----
Total                                     628 kB/s | 45 kB 00:00:00
Running transaction check
Running transaction test
Transaction test succeeded
Running transaction
  Installing : snappy-1.0.5-1.el6.x86_64                1/2
  Installing : snappy-devel-1.0.5-1.el6.x86_64          2/2
  Verifying  : snappy-devel-1.0.5-1.el6.x86_64          1/2
  Verifying  : snappy-1.0.5-1.el6.x86_64                2/2

Installed:
snappy-devel.x86_64 0:1.0.5-1.el6

Dependency Installed:
snappy.x86_64 0:1.0.5-1.el6

Complete!
[ec2-user@ip-10-0-1-13 ~]$
```

Figure 7.11.3

Follow the steps highlighted in Figure 7.11.2 and Figure 7.11.3 to resolve the “snappy-devel” issue.

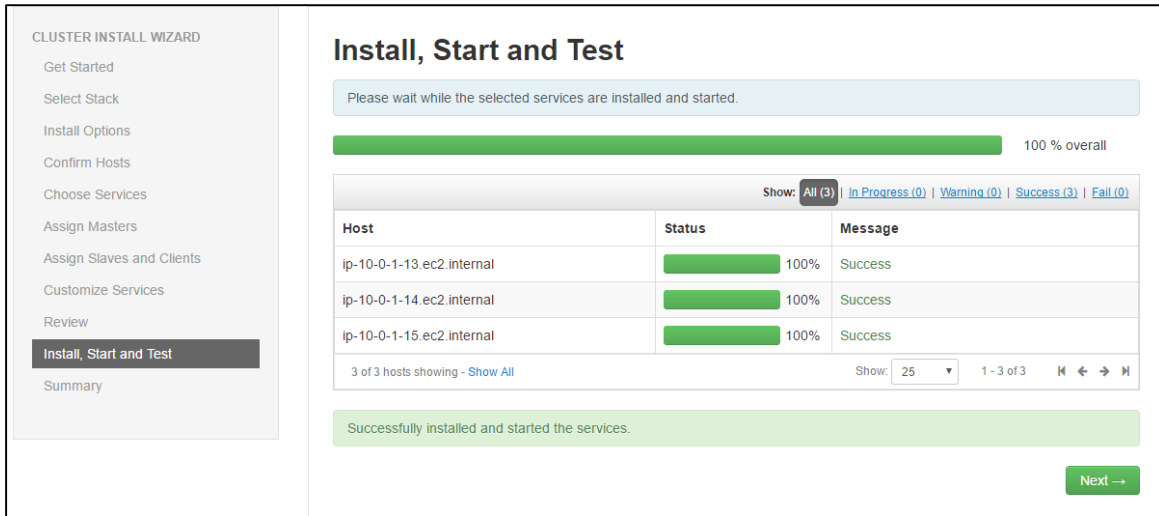


Figure 7.11.4

Once the issue was resolved, I was able successfully install and start all services, as seen Figure 7.11.4.

## 7.12 COMPLETE

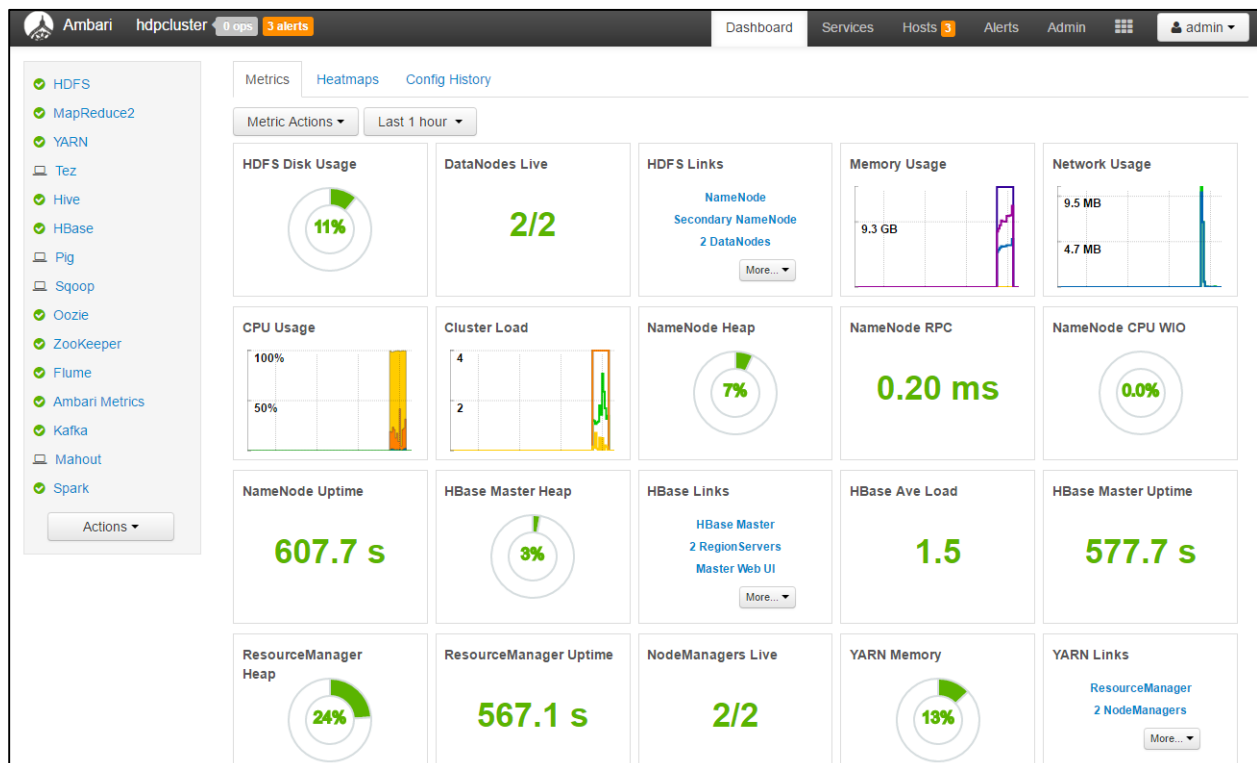


Figure 7.12.1

Your Ambari dashboard will look very similar to Figure 7.12.1, once all your services are up and running.

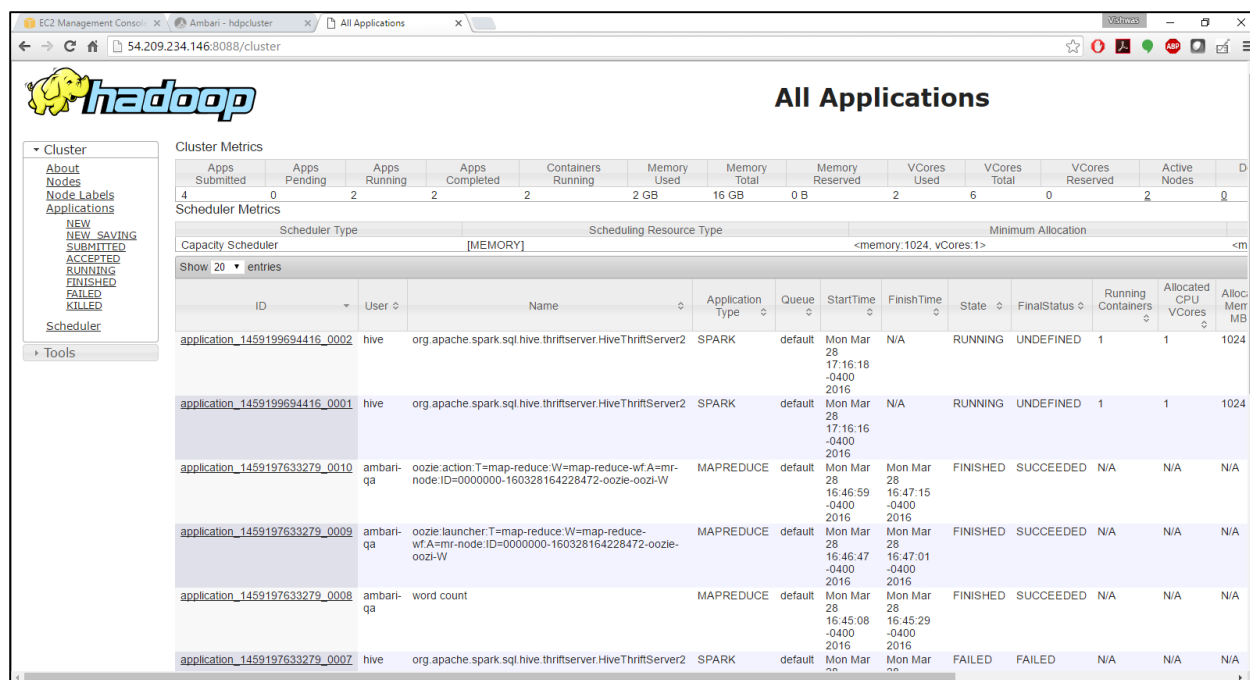


Figure 7.12.2

Figure 7.12.2 – Resource Manager Web UI for the HDP cluster